



PONTIFÍCIA UNIVERSIDADE CATÓLICA DO RIO GRANDE DO SUL
FACULDADE DE INFORMÁTICA
PROGRAMA DE PÓS GRADUAÇÃO EM CIÊNCIA DA COMPUTAÇÃO

MLP and Convolutional Neural Nets

Gabriel Simões



BUSINESS INTELLIGENCE AND
MACHINE LEARNING RESEARCH GROUP

Agenda

- Some words of Machine Learning
- Neural Networks (NN) Overview
 - Multi-layer Perceptron (MLP) training issues
- Convolutional Neural Networks (CNN)
 - CNN training
- Frameworks
 - What kind of DL frameworks we have?
 - Samples

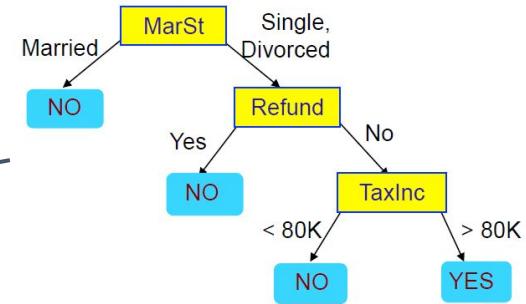
Machine Learning: some simple words

A kind of magic

Machine Learning

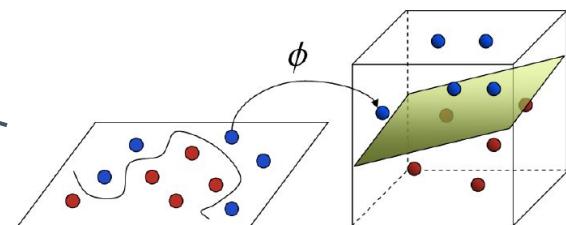
Supervised Learning

- Decision Tree
- SVM
- Naive Bayes
- Neural Network



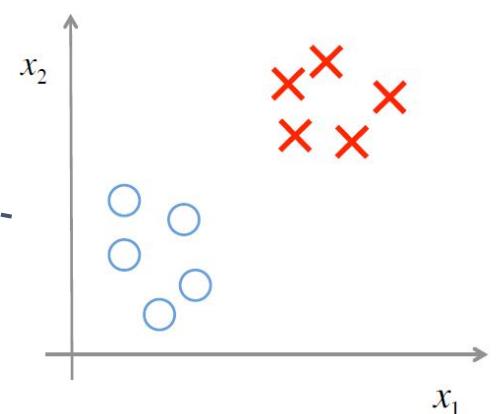
Unsupervised Learning

- EM - Expectation Maximization
- K-means



Reinforcement Learning

- Monte Carlo Methods



Traditional Machine Learning

Is good for:

- Predictions
 - Predict sales
 - Predict prices
 - Pattern recognition
- Clustering
 - Customer profile
 - Anomaly detection

Not so good for **high dimensionality** problems!

Such image, video, audio...



Deep Learning

“A family of methods that use **deep architectures** to learn **high-level features abstractions**.”

Kevin Duh (Nara Institute of Science and Technology)

Deep Learning

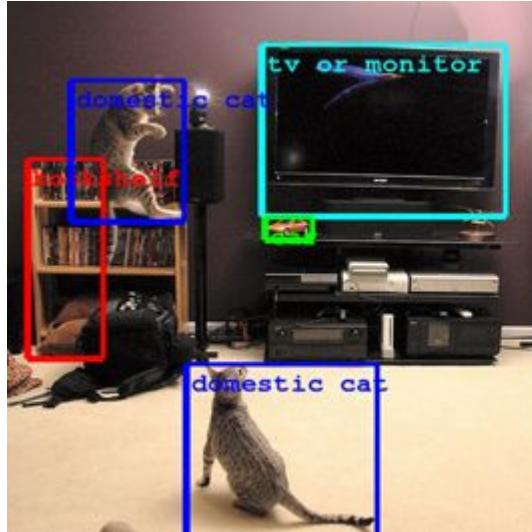
“A family of methods that use **deep architectures** to learn **high-level features abstractions**.”

Kevin Duh (Nara Institute of Science and Technology)

Deep Learning

Is good for:

Image



Video



Audio



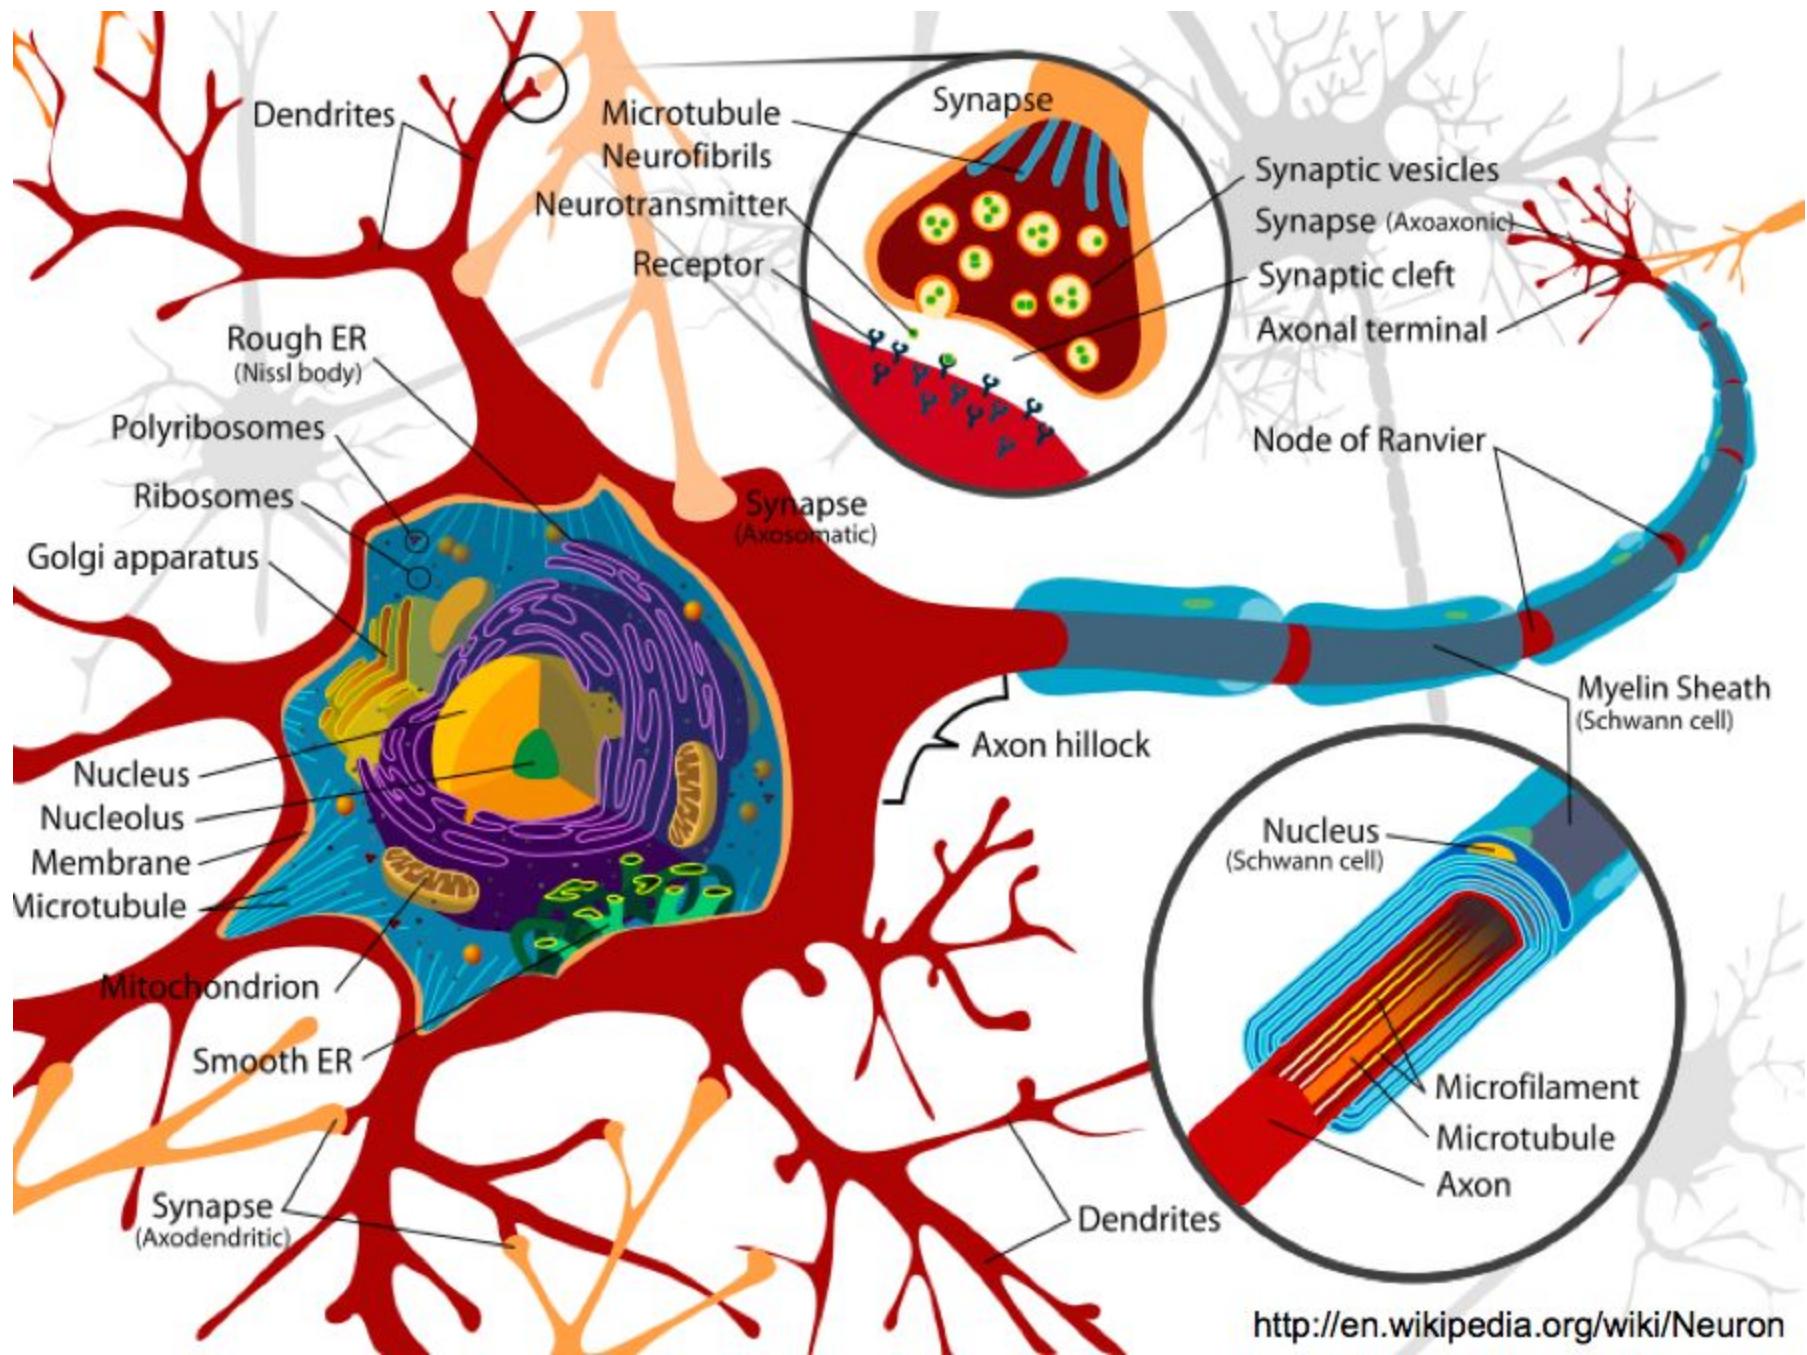
Introduction to Neural Networks

How the magic happens

Neural Networks

A neural network is a massively parallel distributed processor made up of simple processing units that has a natural propensity for storing experiential knowledge and making it available for use. It resembles the brain in two respects:

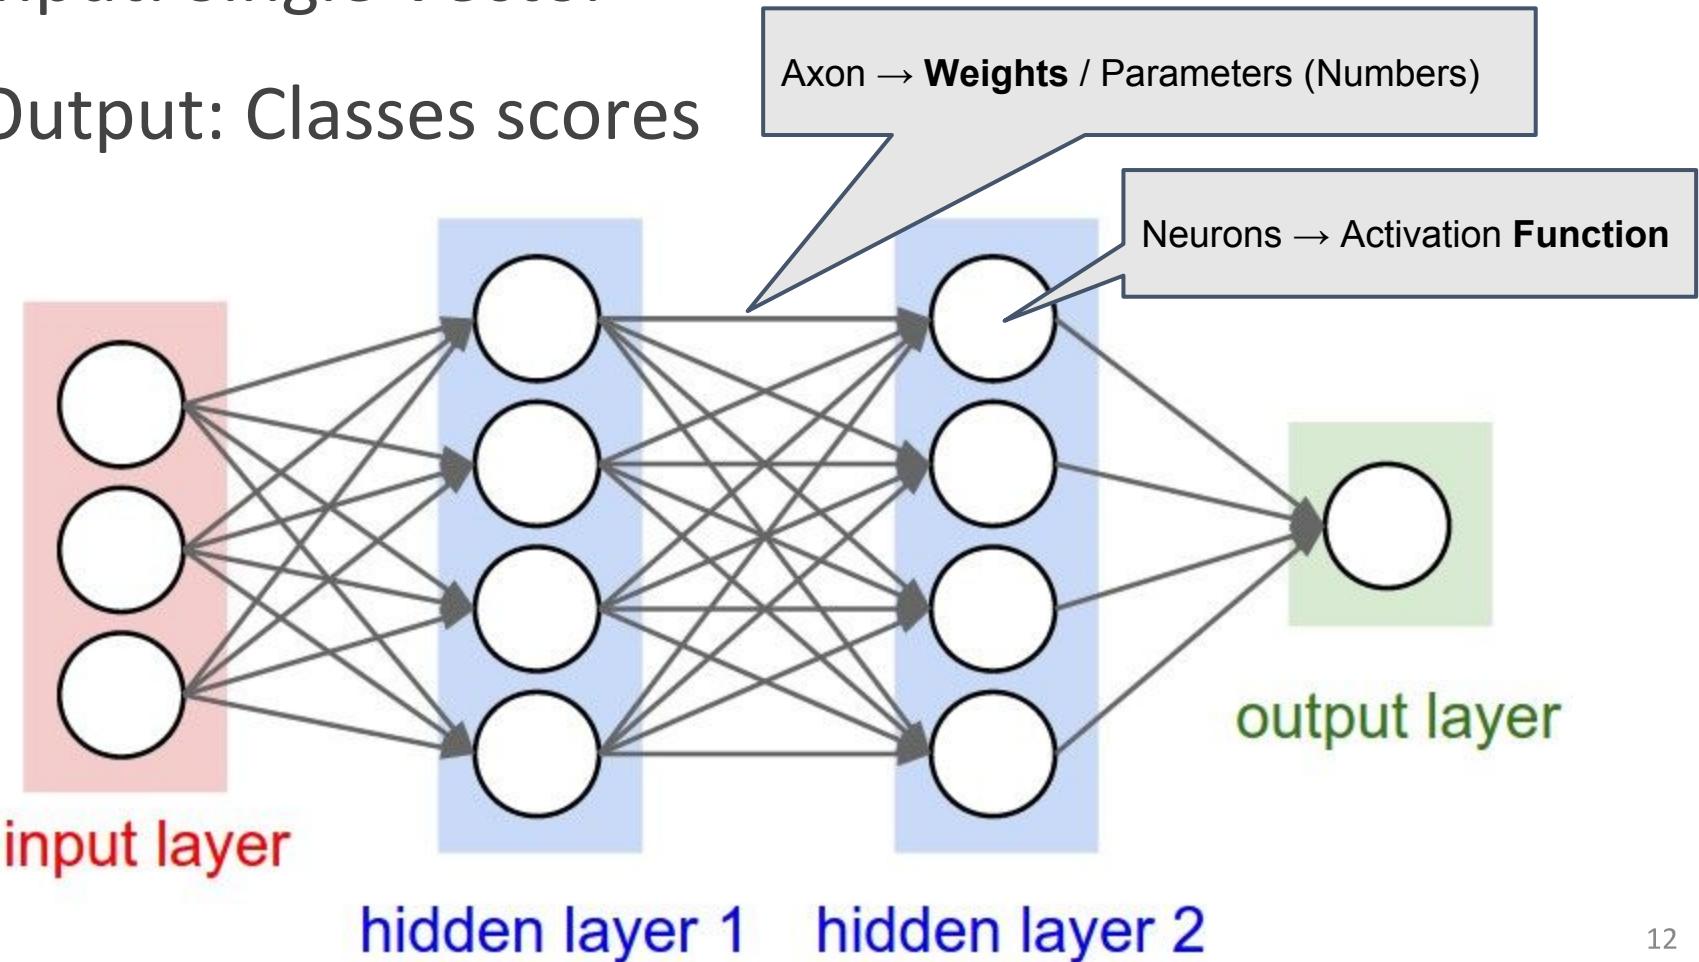
1. Knowledge is acquired by the network from its environment through a learning process.
2. Interneuron connection strengths, known as synaptic weights, are used to store the acquired knowledge.



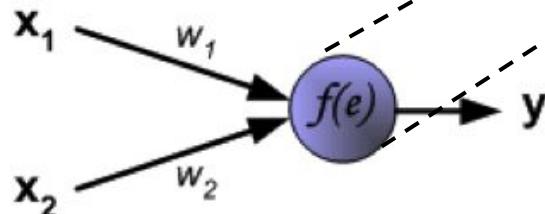
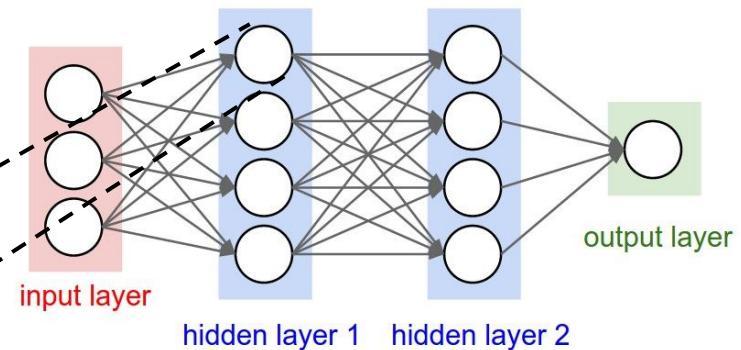
Neural Networks

Input: Single Vector

Output: Classes scores

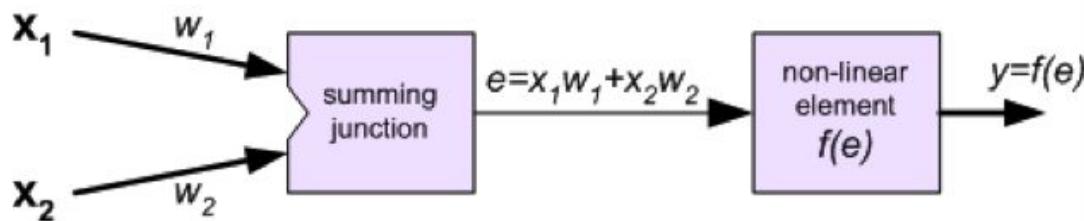


Neural Networks



$$f(x) = \frac{1}{1 + e^{-x}}$$

Activation Function

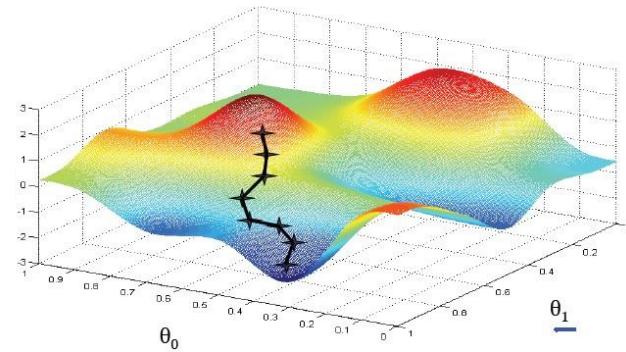


Backpropagation

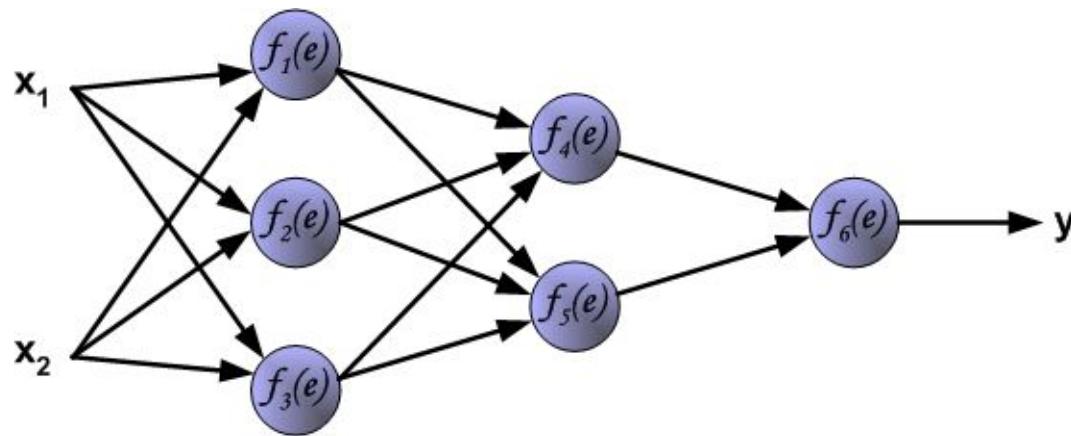
A technique for implementing gradient descent in weight space for a multilayer perceptron.

Training steps:

- For each training instance
 - Input Data
 - Feed-Forward
 - Output error calculation
 - Error Backpropagation
 - Weights adjust

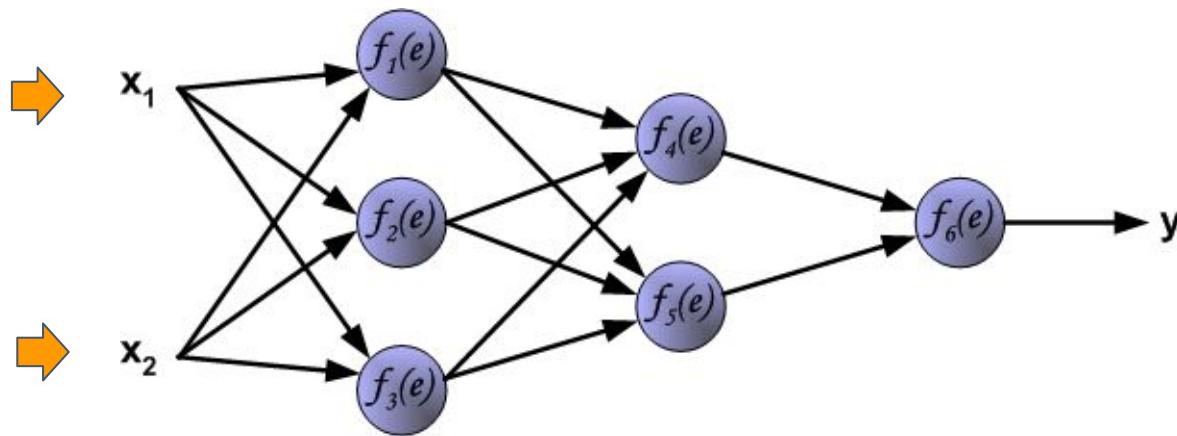


Backpropagation



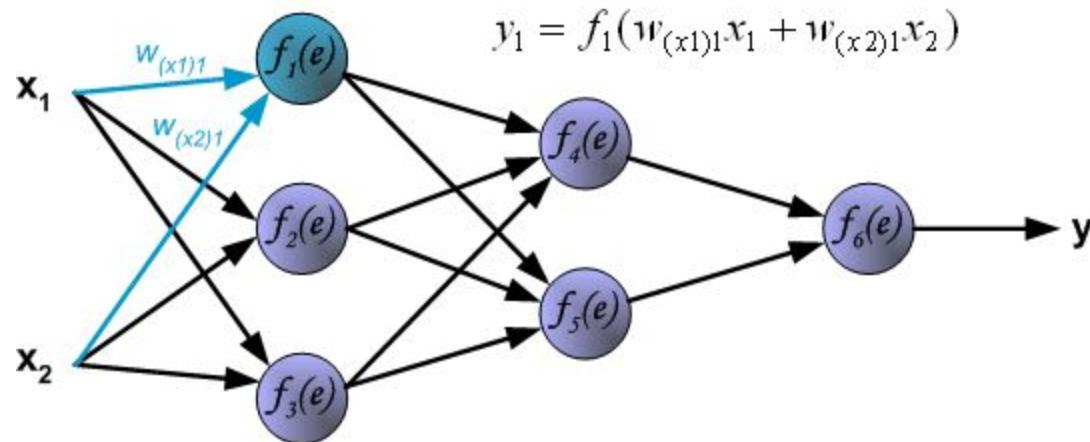
Backpropagation

Input Data



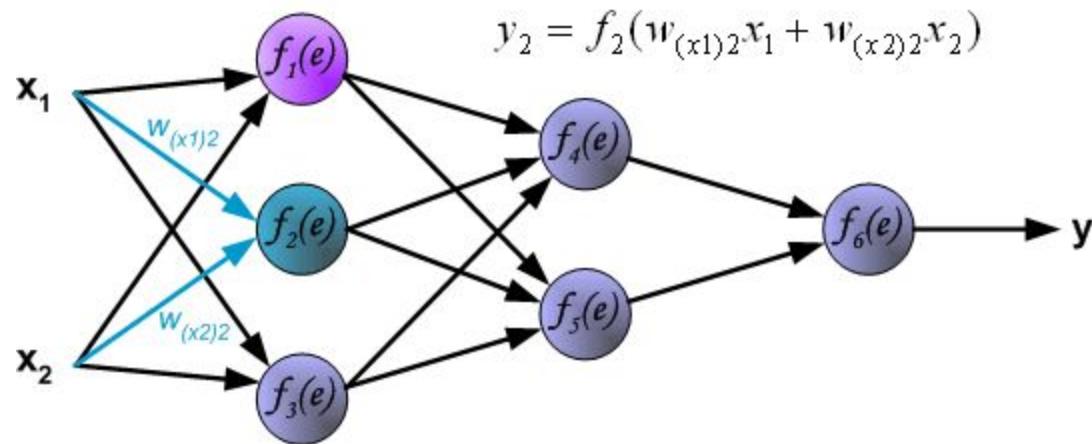
Backpropagation

Feed-forward



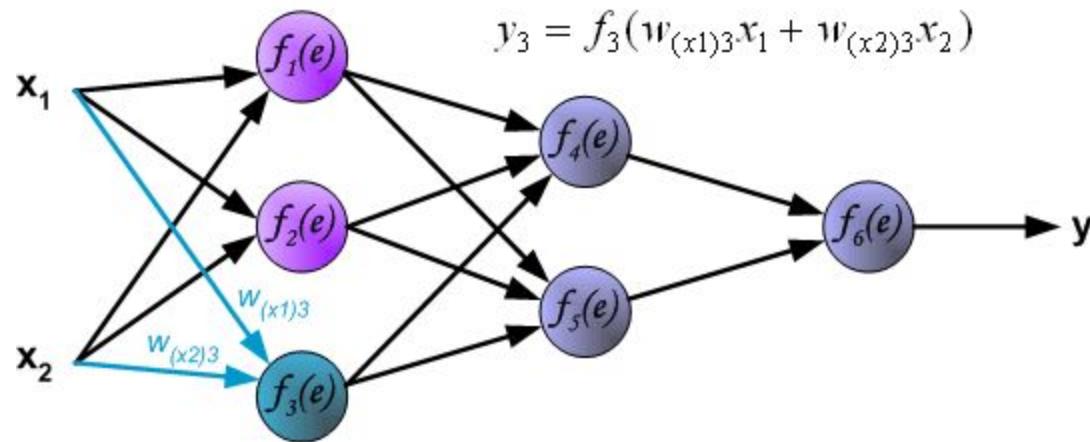
Backpropagation

Feed-forward



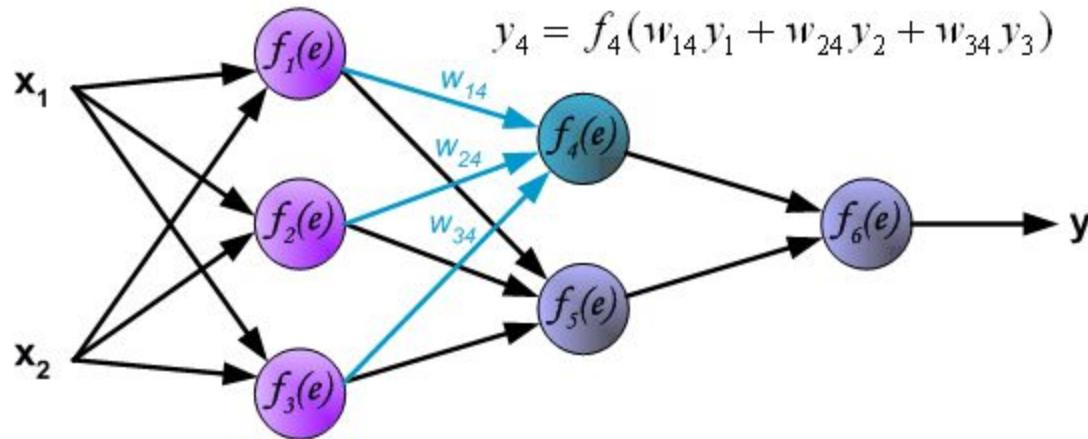
Backpropagation

Feed-forward



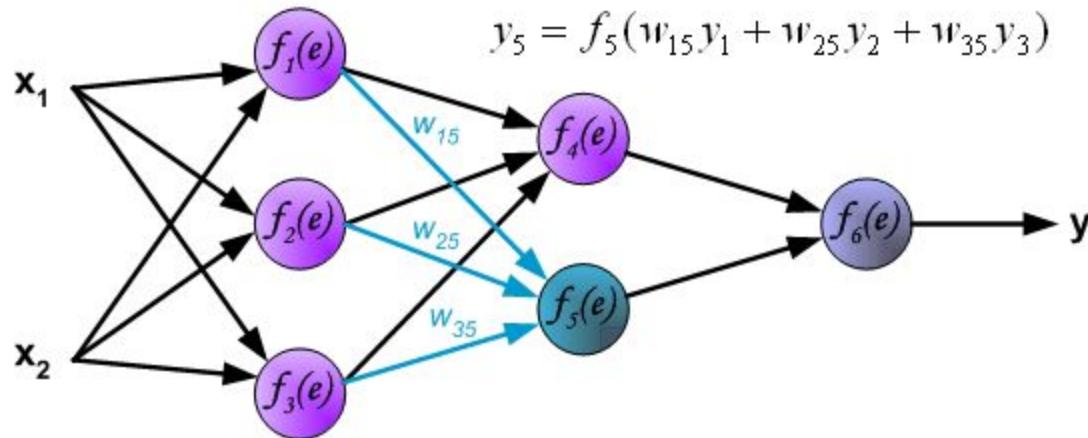
Backpropagation

Feed-forward



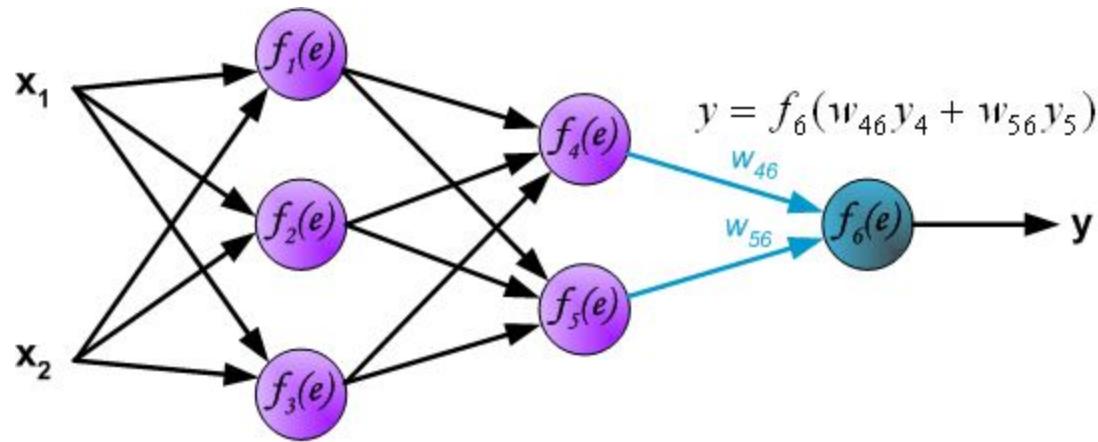
Backpropagation

Feed-forward



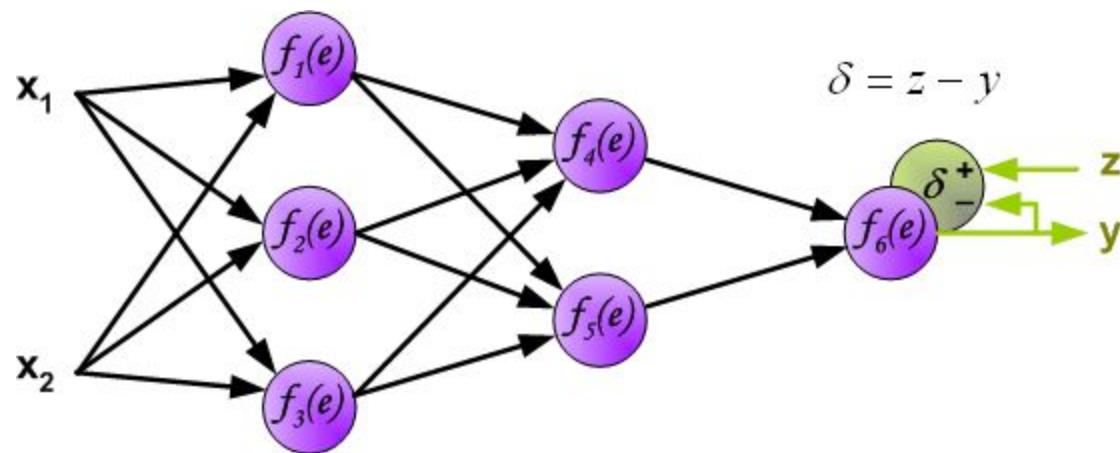
Backpropagation

Feed-forward



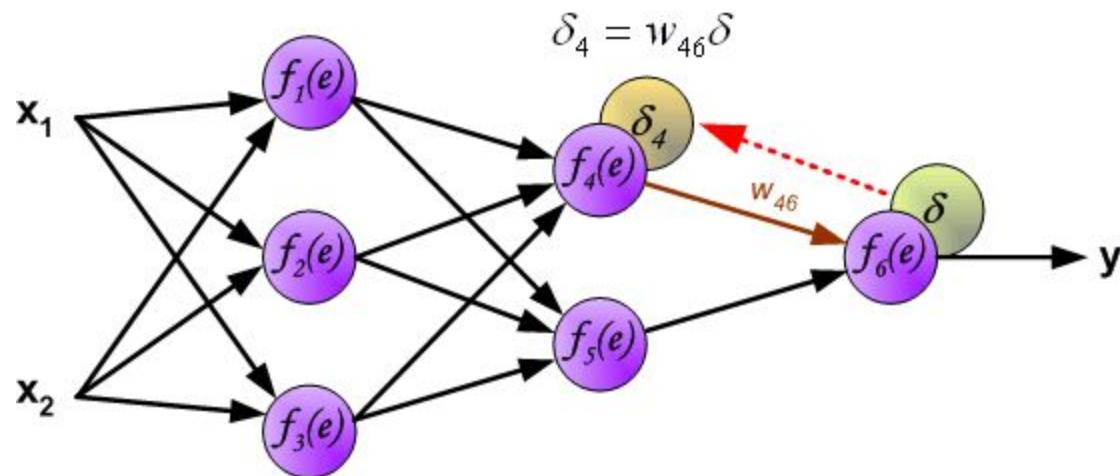
Backpropagation

Output error calculation



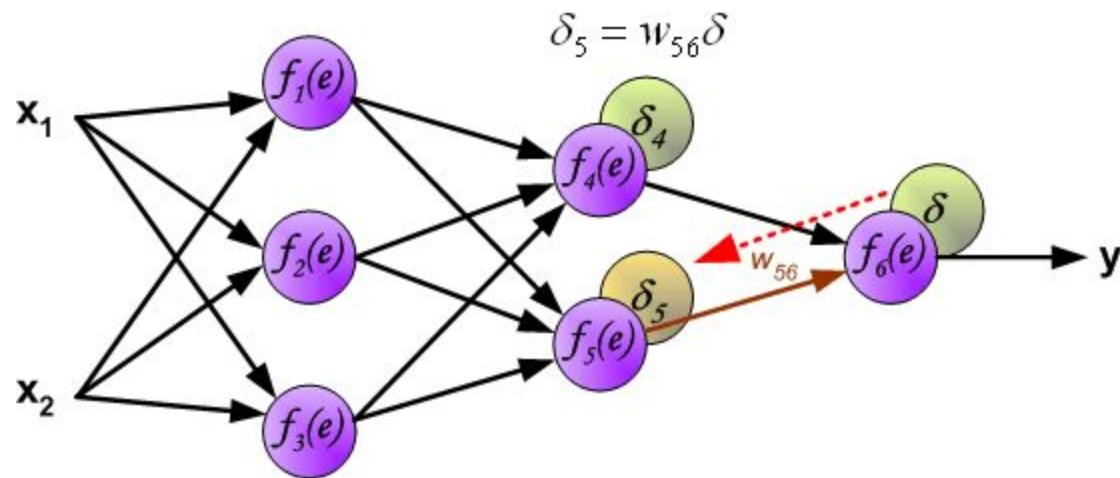
Backpropagation

Backpropagate error



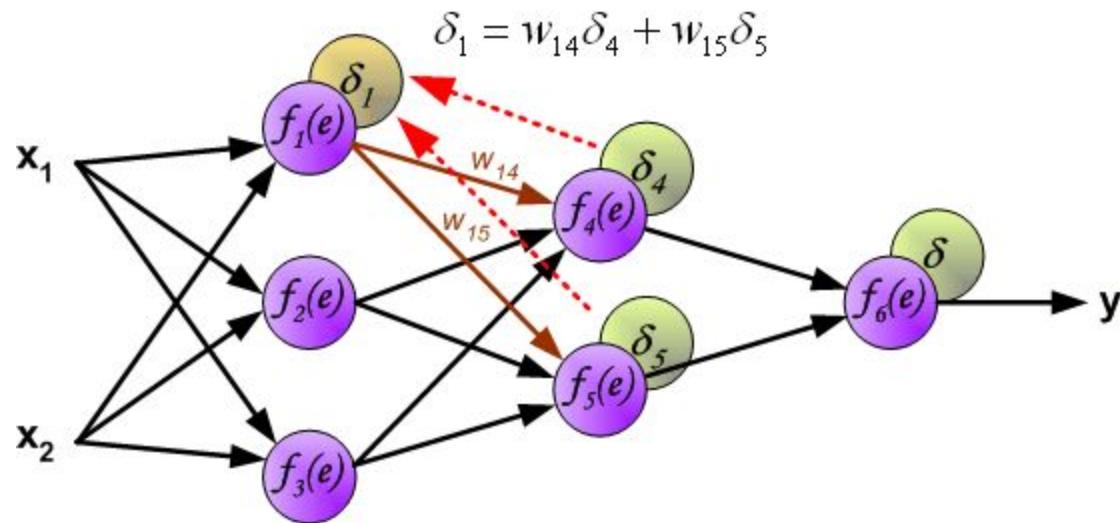
Backpropagation

Backpropagate error



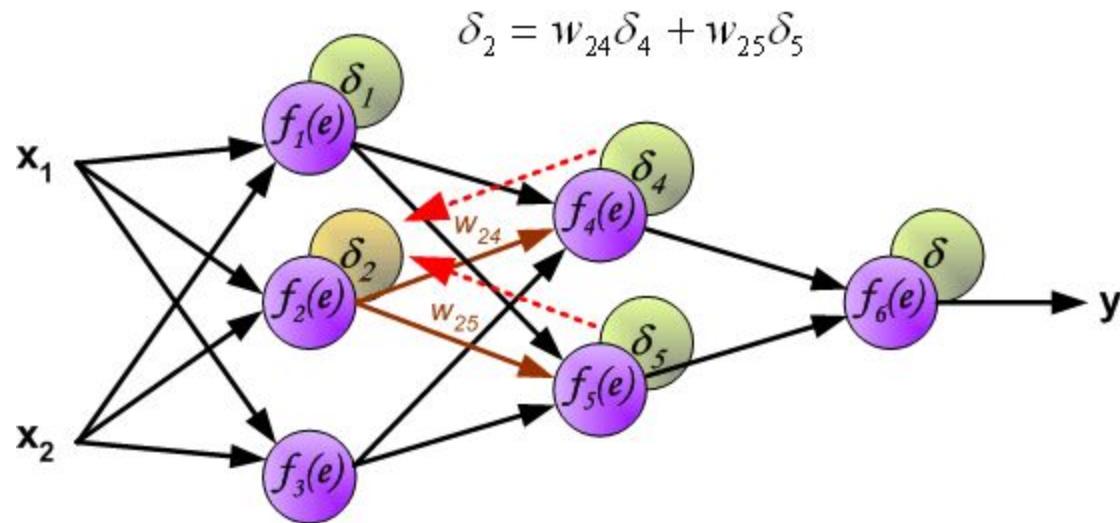
Backpropagation

Backpropagate error



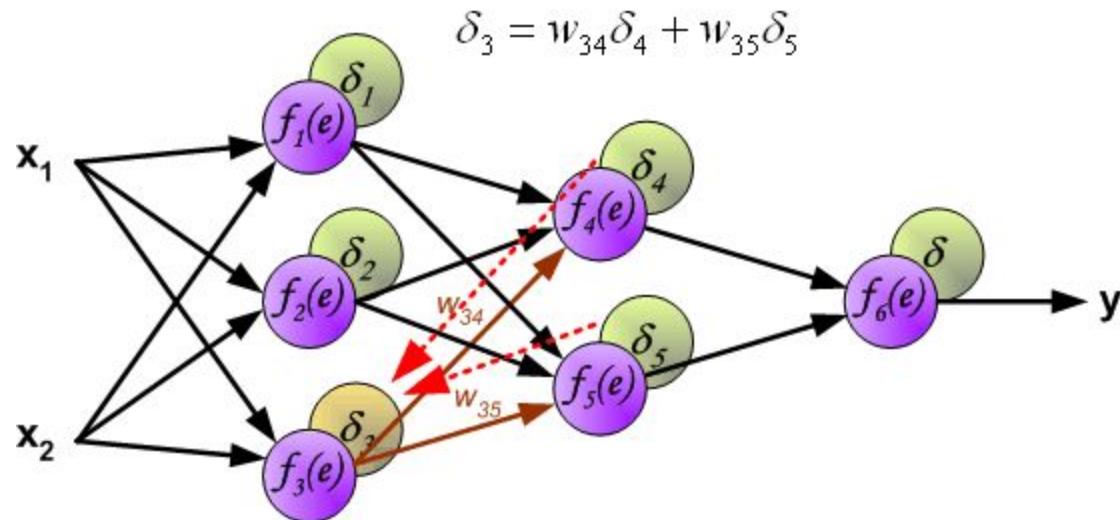
Backpropagation

Backpropagate error



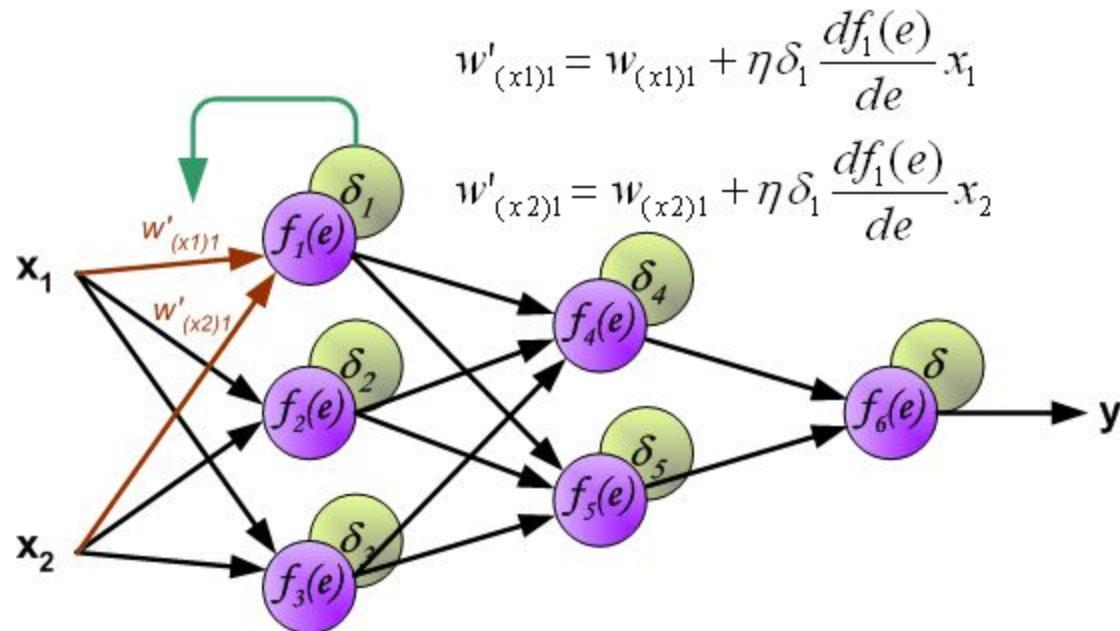
Backpropagation

Backpropagate error



Backpropagation

Weights update

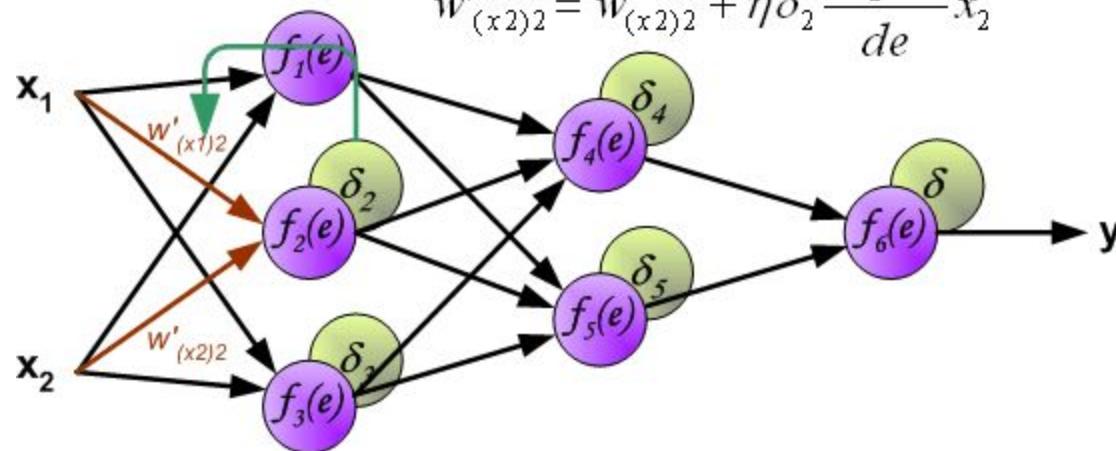


Backpropagation

Weights update

$$w'_{(x1)2} = w_{(x1)2} + \eta \delta_2 \frac{df_2(e)}{de} x_1$$

$$w'_{(x2)2} = w_{(x2)2} + \eta \delta_2 \frac{df_2(e)}{de} x_2$$

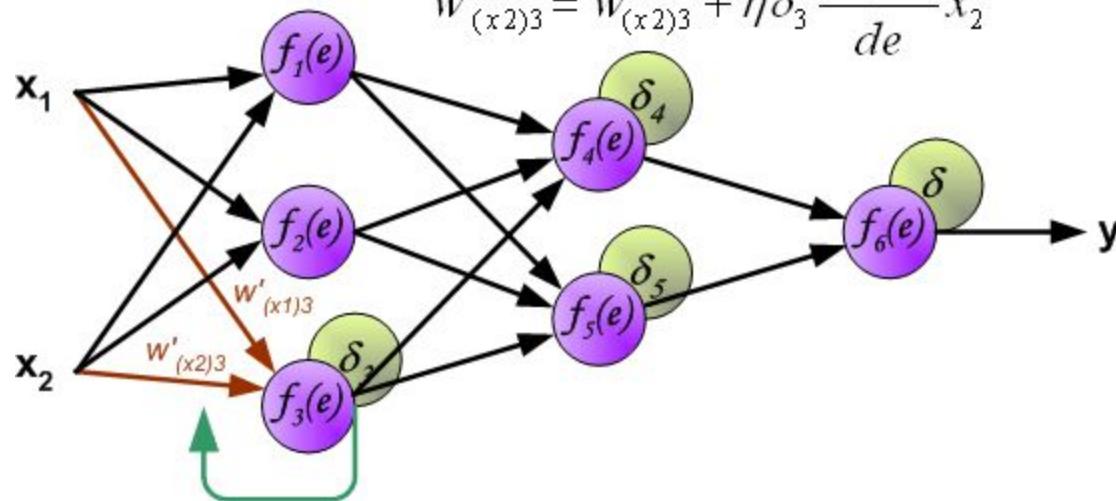


Backpropagation

Weights update

$$w'_{(x1)3} = w_{(x1)3} + \eta \delta_3 \frac{df_3(e)}{de} x_1$$

$$w'_{(x2)3} = w_{(x2)3} + \eta \delta_3 \frac{df_3(e)}{de} x_2$$

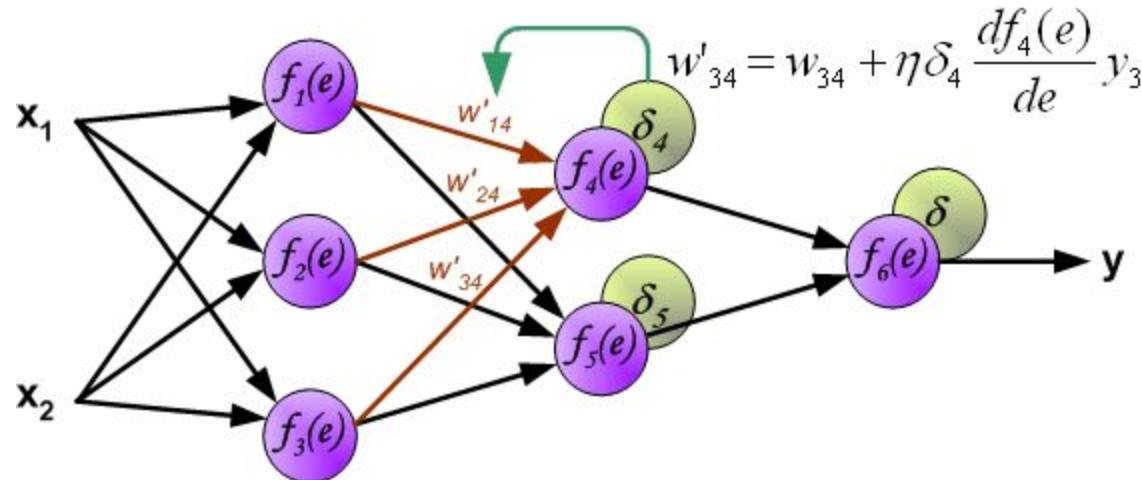


Backpropagation

Weights update

$$w'_{14} = w_{14} + \eta \delta_4 \frac{df_4(e)}{de} y_1$$

$$w'_{24} = w_{24} + \eta \delta_4 \frac{df_4(e)}{de} y_2$$



http://galaxy.agh.edu.pl/~vlsi/AI/backp_t_en/backprop.html

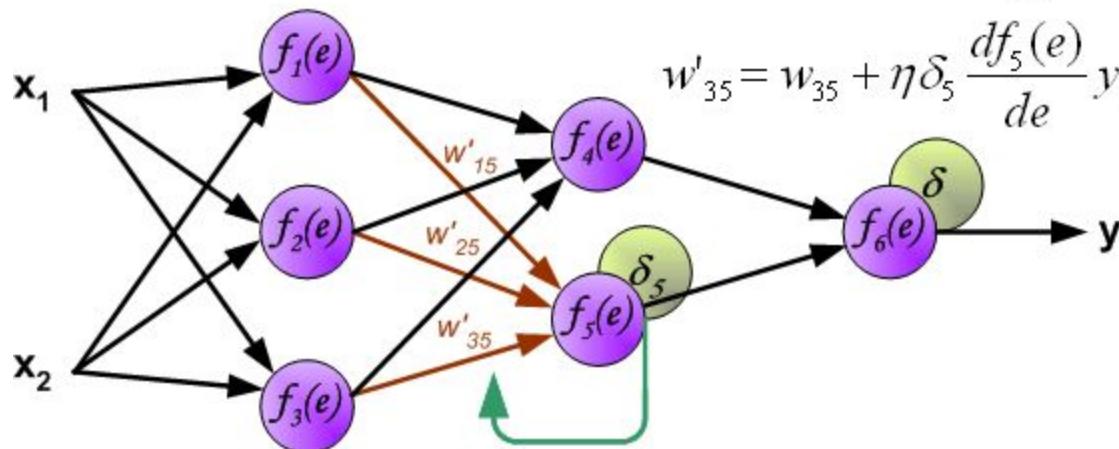
Backpropagation

Weights update

$$w'_{15} = w_{15} + \eta \delta_5 \frac{df_5(e)}{de} y_1$$

$$w'_{25} = w_{25} + \eta \delta_5 \frac{df_5(e)}{de} y_2$$

$$w'_{35} = w_{35} + \eta \delta_5 \frac{df_5(e)}{de} y_3$$



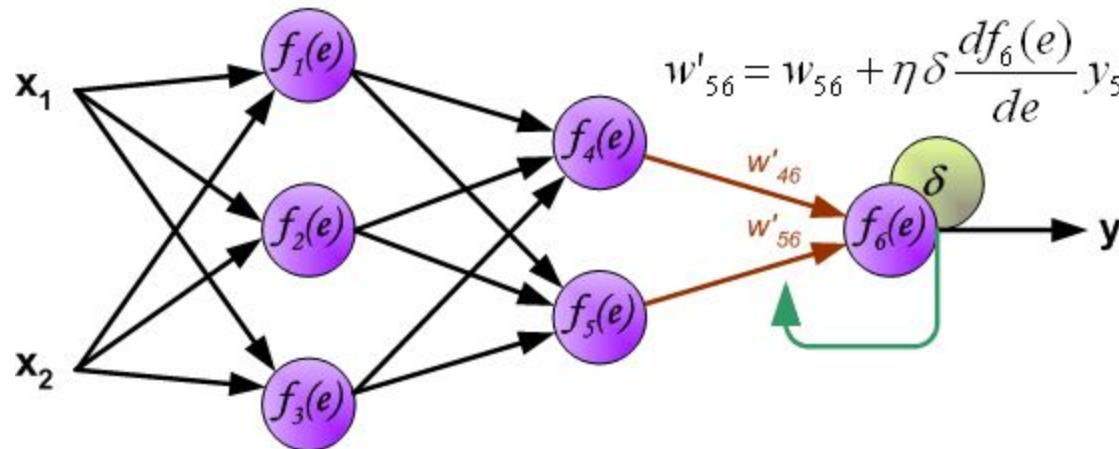
http://galaxy.agh.edu.pl/~vlsi/AI/backp_t_en/backprop.html

Backpropagation

Weights update

$$w'_{46} = w_{46} + \eta \delta \frac{df_6(e)}{de} y_4$$

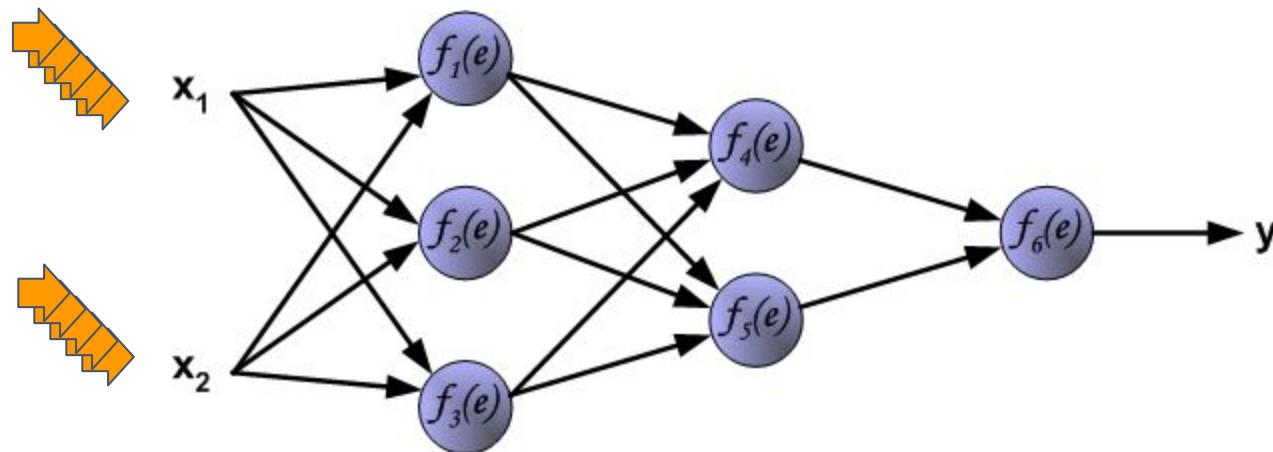
$$w'_{56} = w_{56} + \eta \delta \frac{df_6(e)}{de} y_5$$



http://galaxy.agh.edu.pl/~vlsi/AI/backp_t_en/backprop.html

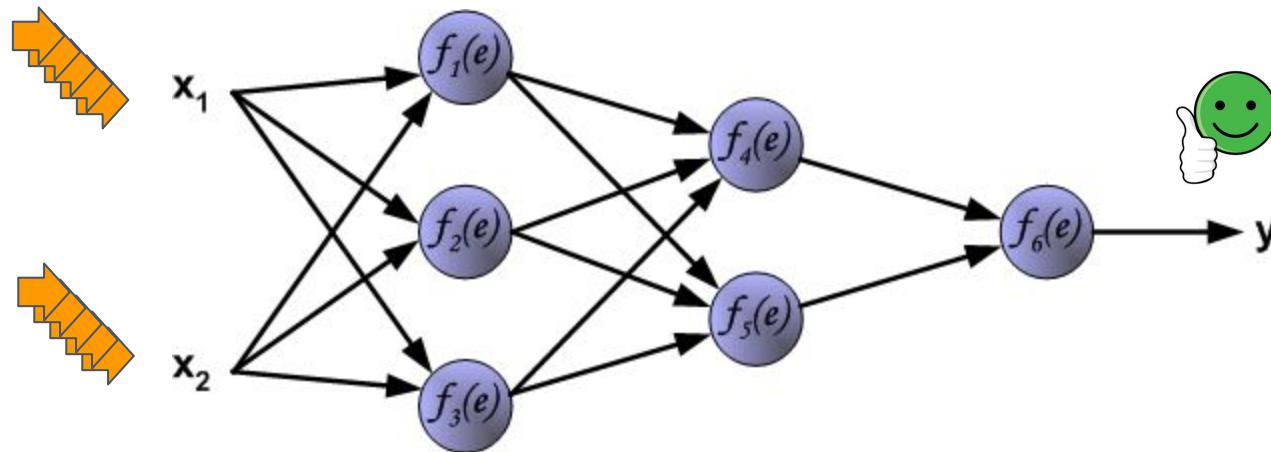
Backpropagation

Repeat the process for every training instance...



Backpropagation

Repeat the process for every training instance...



... n times or until the output error be acceptable.

Neural Networks

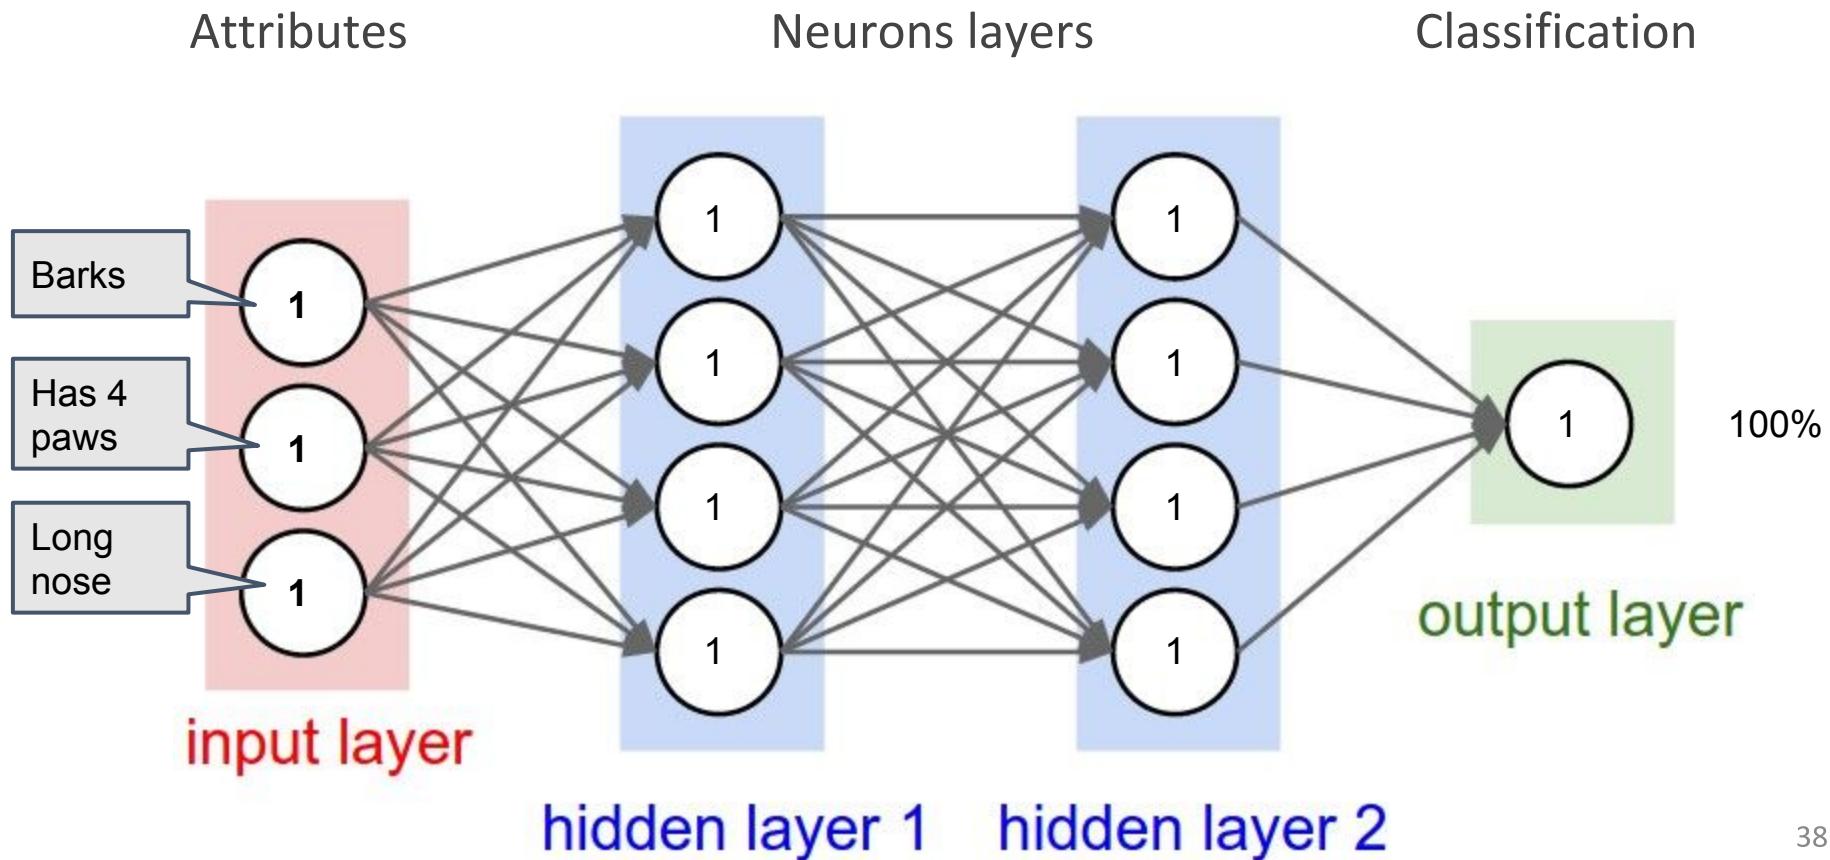
$$\frac{1}{1 + e^{-x}}$$

Activation Function Derivative

$$\begin{aligned}
 \frac{\partial \sigma}{\partial x} &= \frac{\partial}{\partial x} \left[\frac{1}{1 + e^{-x}} \right] \\
 &= \frac{\partial}{\partial x} (1 + e^{-x})^{-1} \\
 &= -(1 + e^{-x})^{-2} (-e^{-x}) \\
 &= \frac{e^{-x}}{(1 + e^{-x})^2} \\
 &= \frac{1}{1 + e^{-x}} \cdot \frac{e^{-x}}{1 + e^{-x}} \\
 &= \frac{1}{1 + e^{-x}} \cdot \frac{(1 + e^{-x}) - 1}{1 + e^{-x}} \\
 &= \frac{1}{1 + e^{-x}} \cdot \left(1 - \frac{1}{1 + e^{-x}} \right) \\
 &= \sigma(x) \cdot (1 - \sigma(x))
 \end{aligned}$$

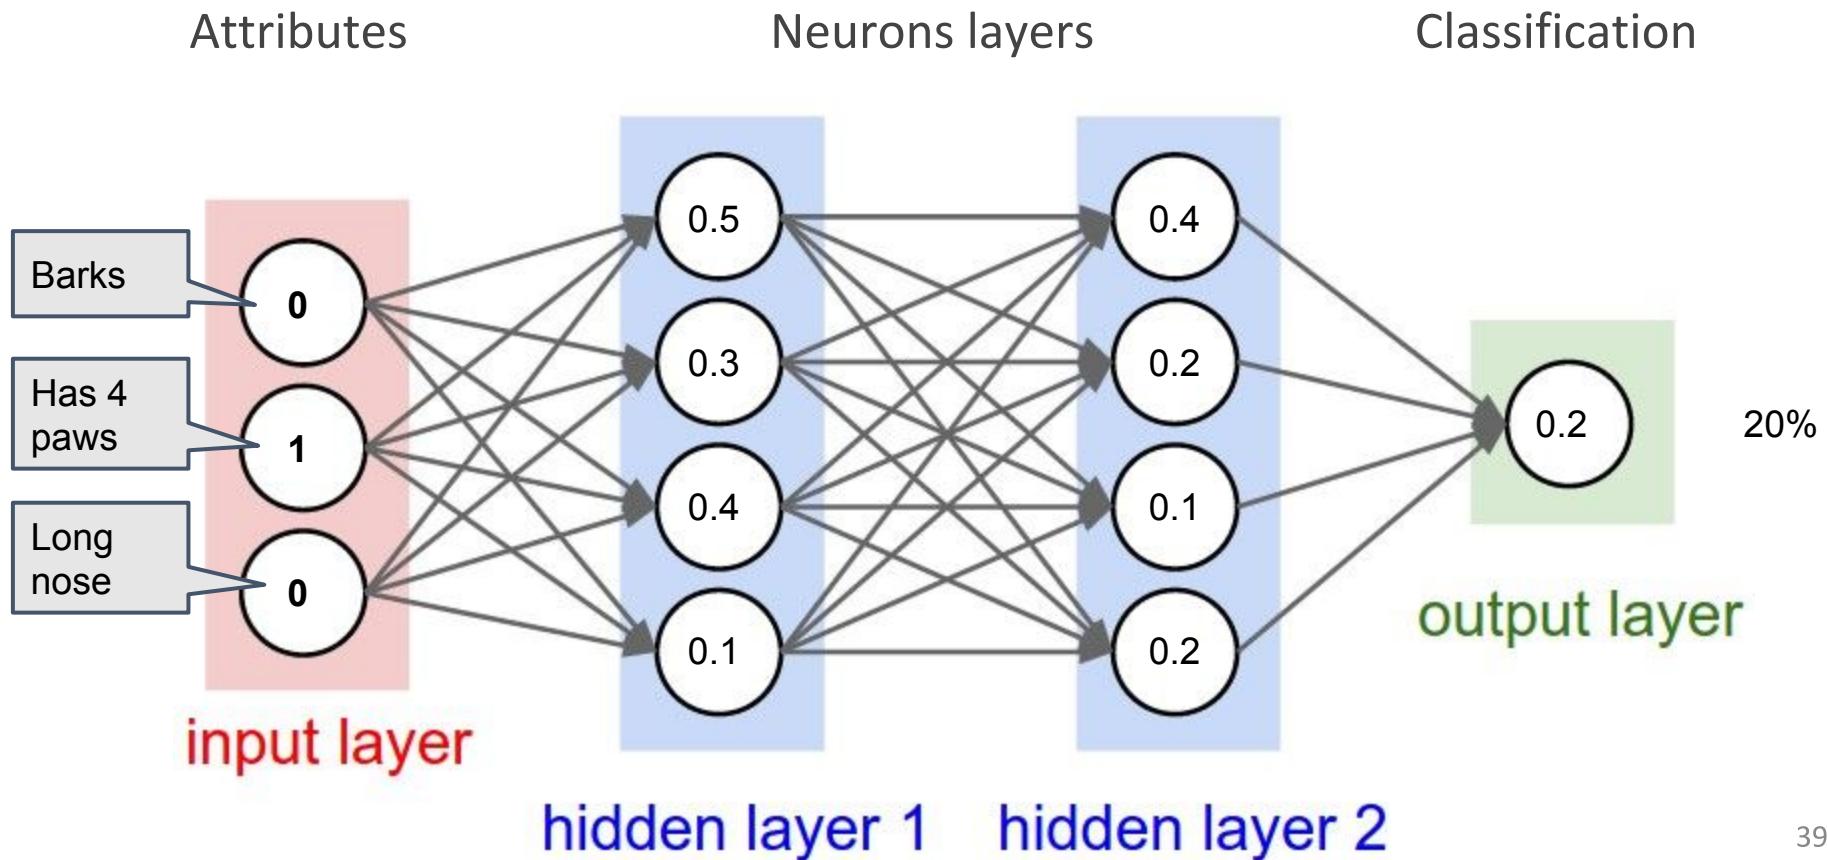
Neural Networks

Toy network for detecting dogs:



Neural Networks

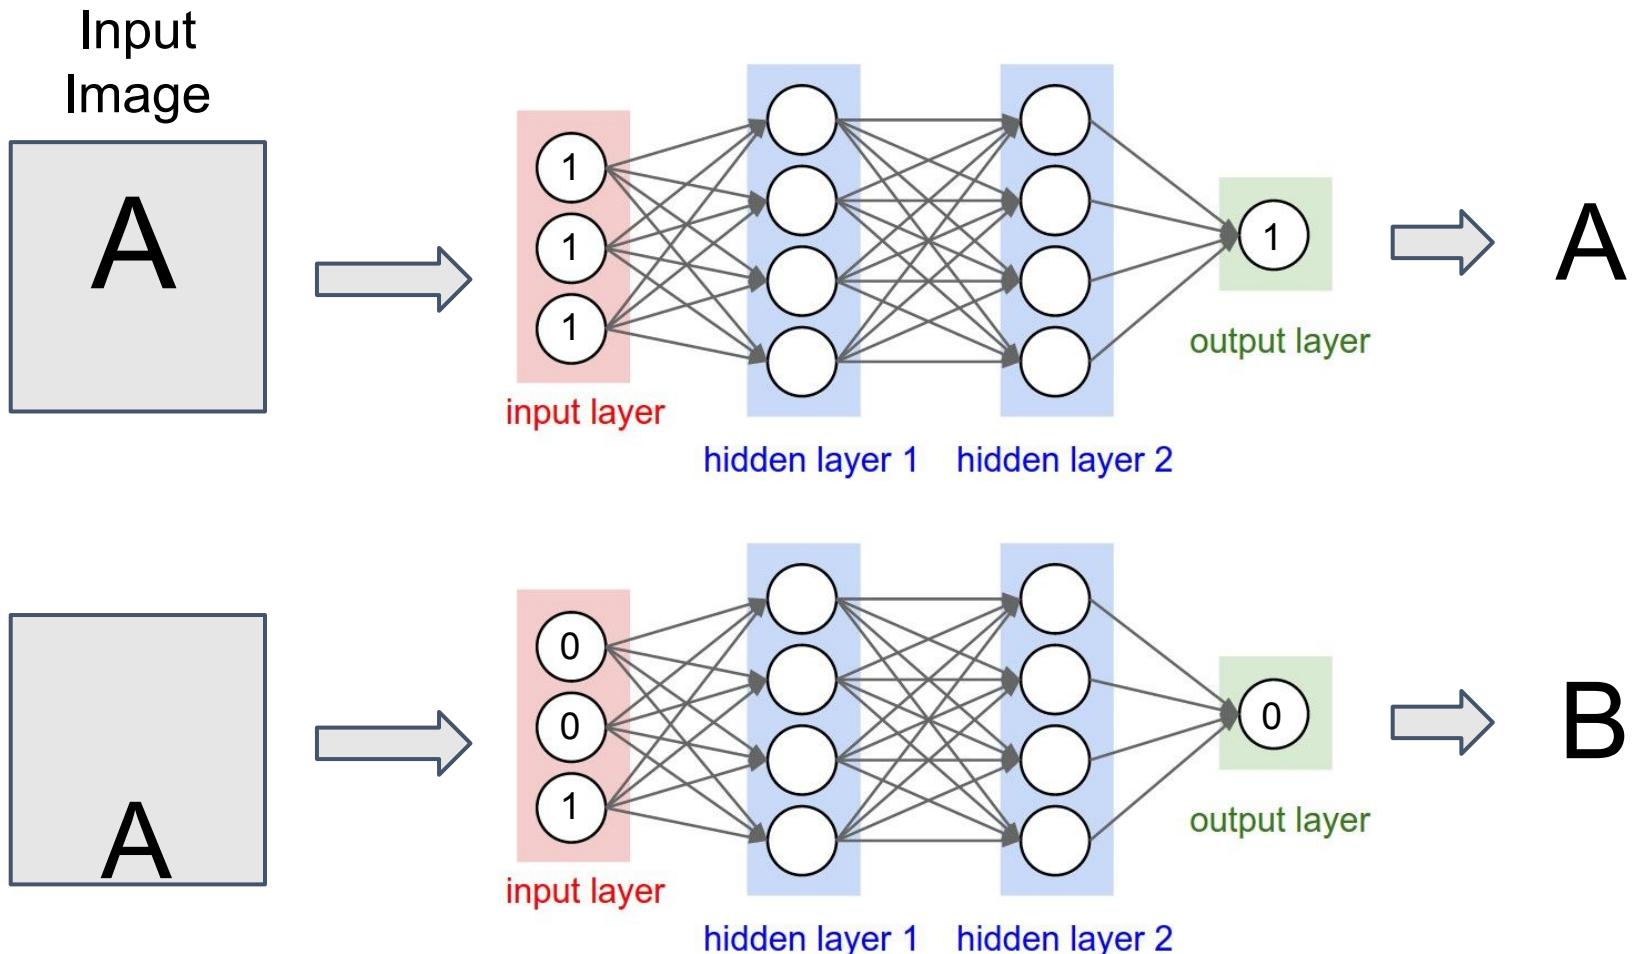
Toy network for detecting dogs:



The image classification challenge

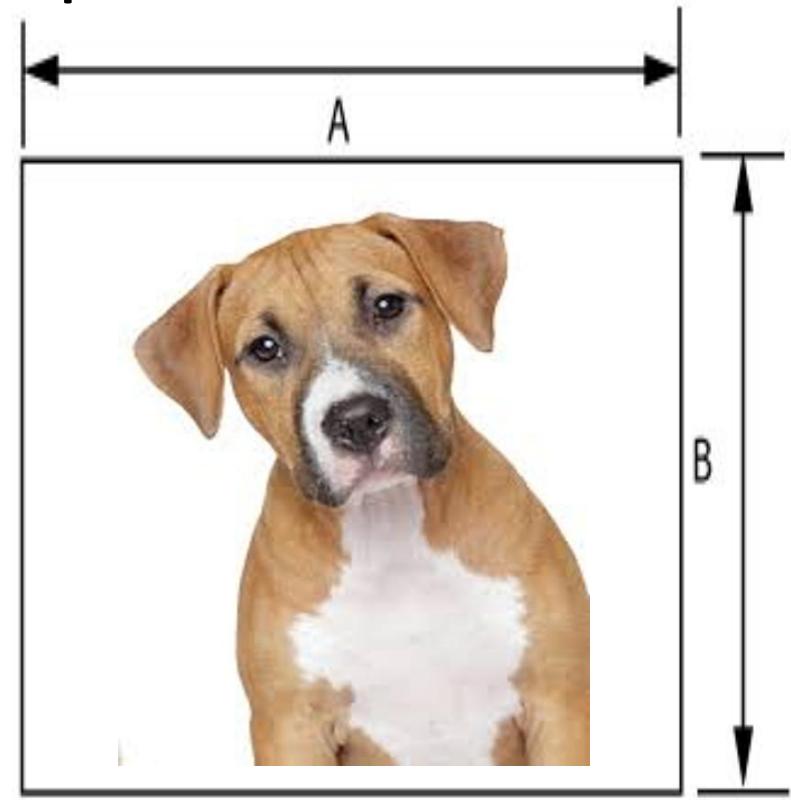
- Image 32 height x 32 width x 3 channels (RGB)
 - 3072 weights (connections)
- Image 200 height x 200 width x 3 channels (RGB)
 - 120,000 weights
- This full connectivity is wasteful and the huge number of parameters would quickly lead to overfitting.
- MLP are not sensitive to input patterns invariance.

Neural Network Invariance Problem



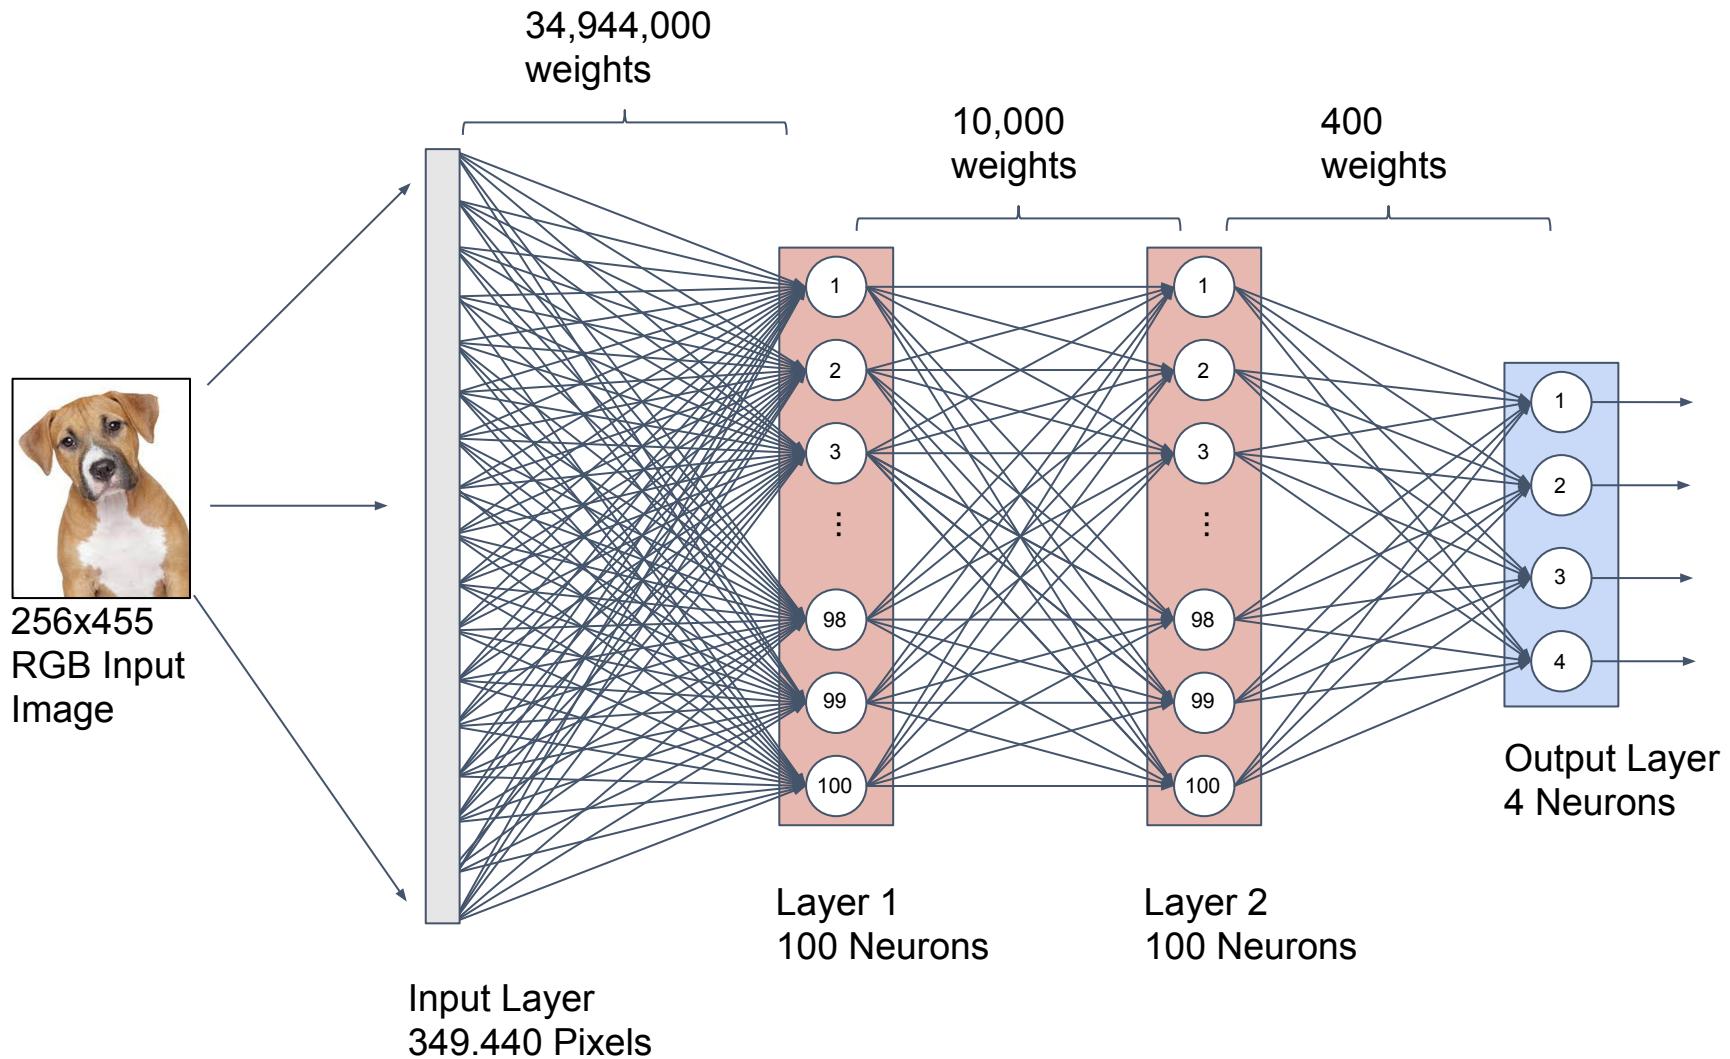
The image classification problem

Dimension

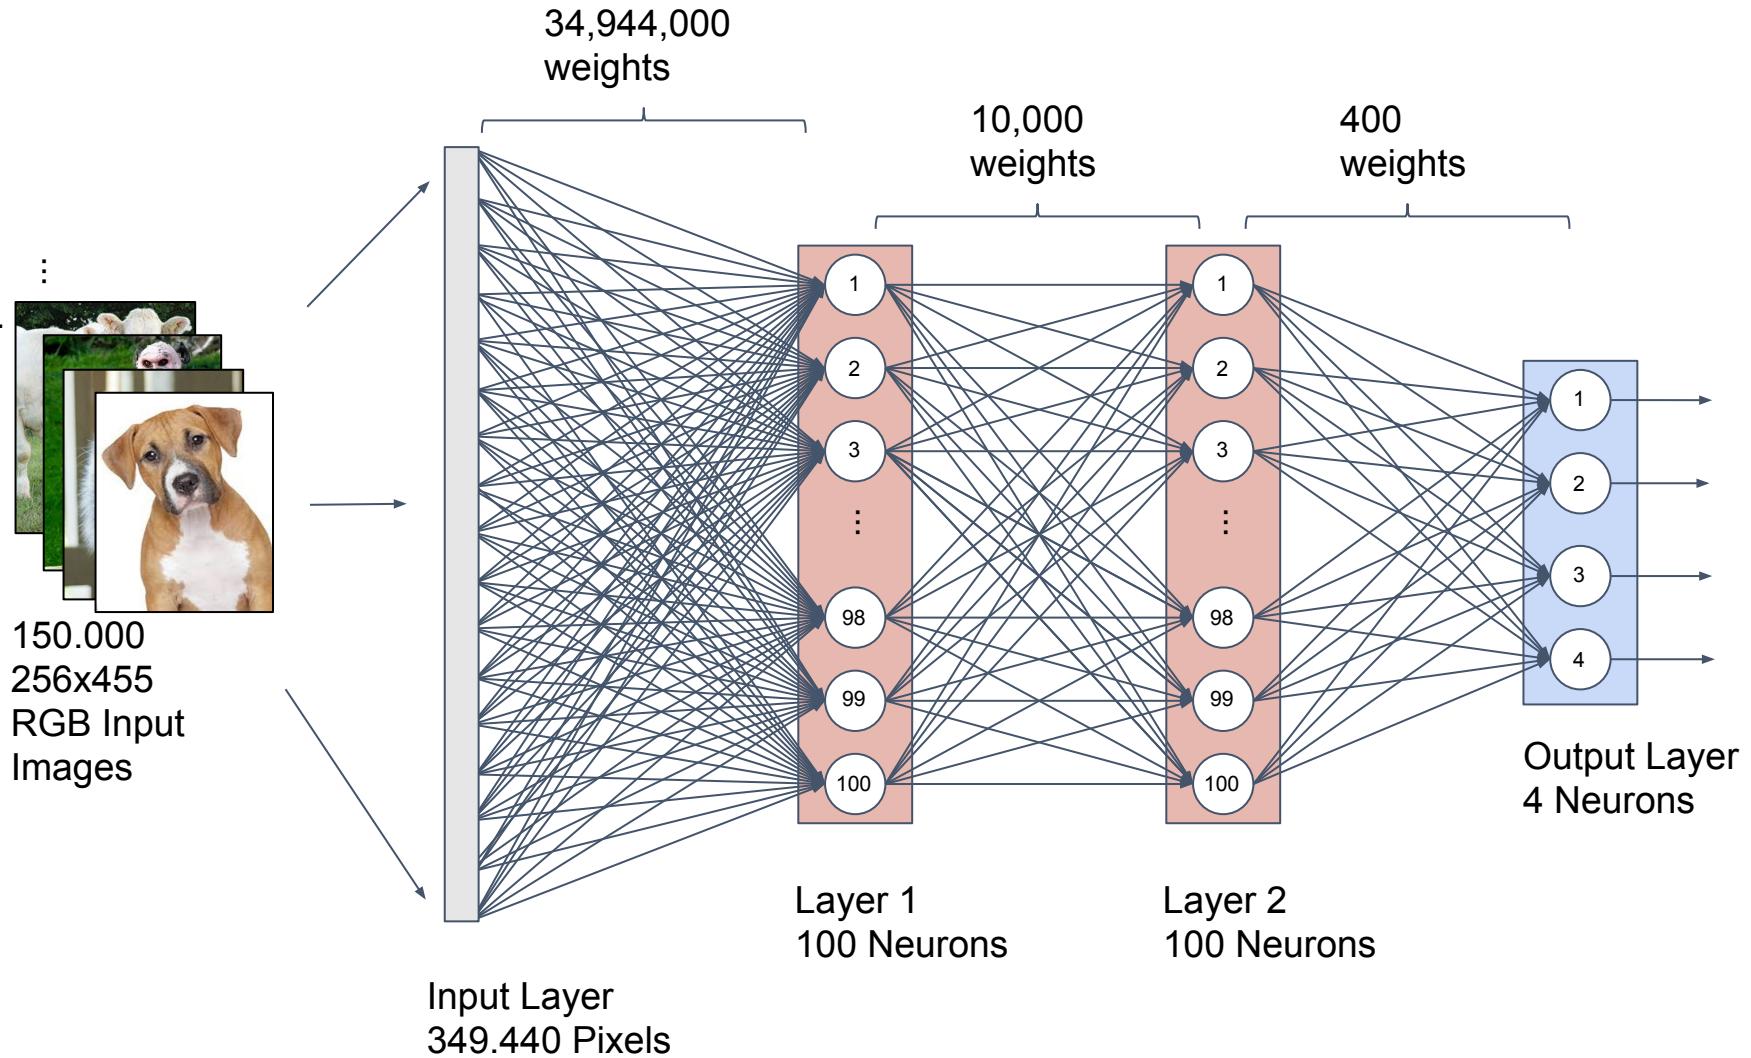


For a $256 \times 455 \times 3$ image we have 349,440 inputs.
How to train a neural net with an input layer like that?

Learning a 256x455 RGB Dog and...



... a Dog in 150k 256x455 RGB Images!



The image classification problem

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation



Convolutional Neural Network

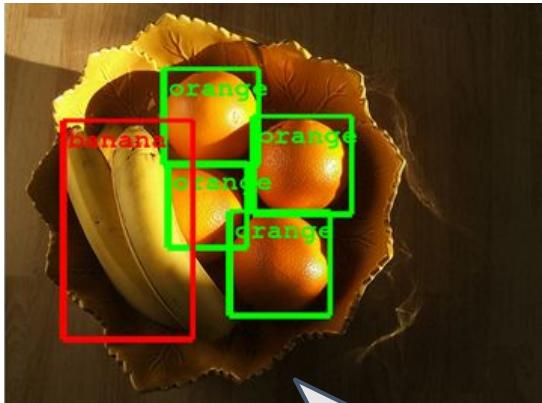
When the magic happens again!

The Convolutional Neural Network

Convolutional neural networks (CNN) are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers.

Applications

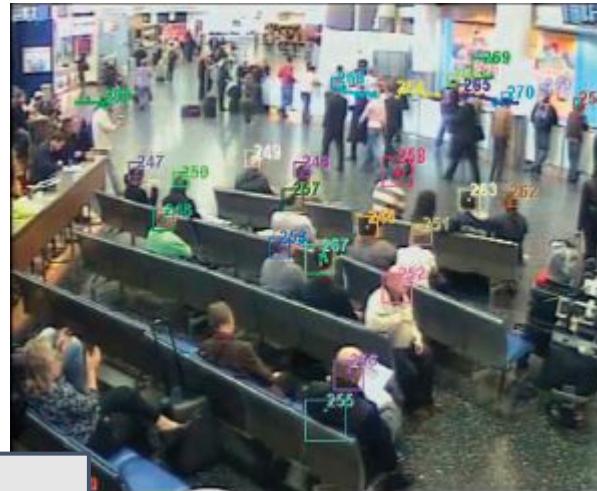
Image Recognition



Scalable Object Detection using Deep Neural Networks
Erhan D. et al (2014)

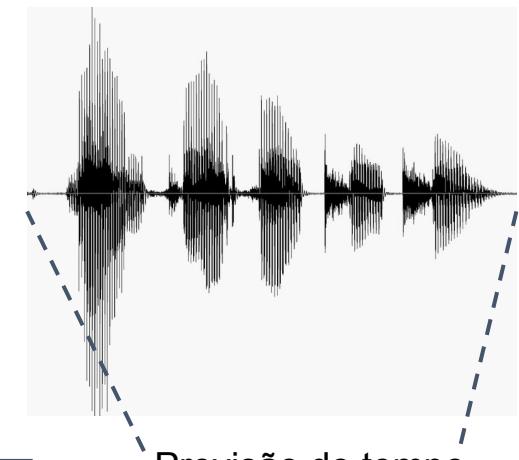
ImageNet Classification with Deep Convolutional Neural Networks.
Alex Krizhevsky et al (2012).

Video Analysis



3D Convolutional Neural Networks for Human Action Recognition
Shuiwang Ji et al (2013)

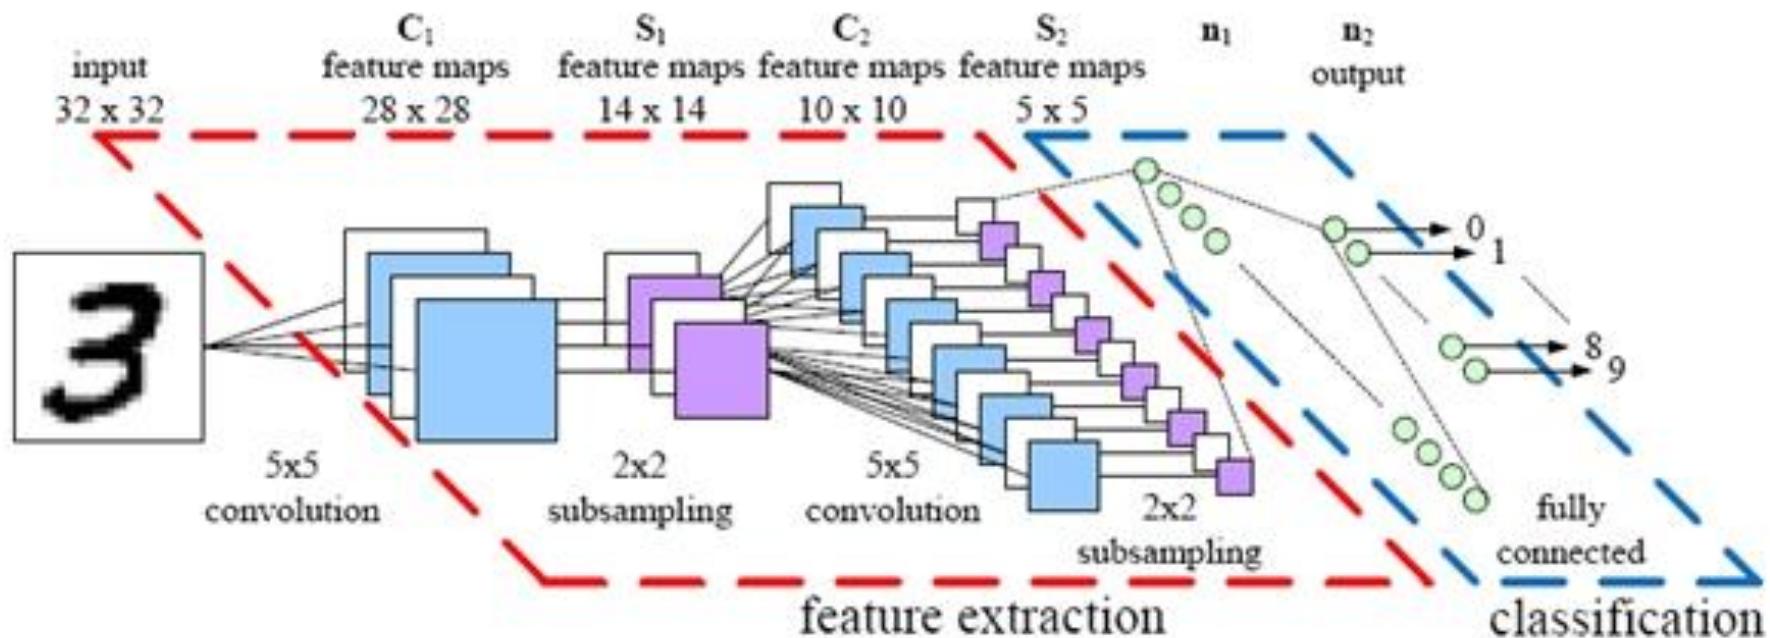
Natural Language Processing



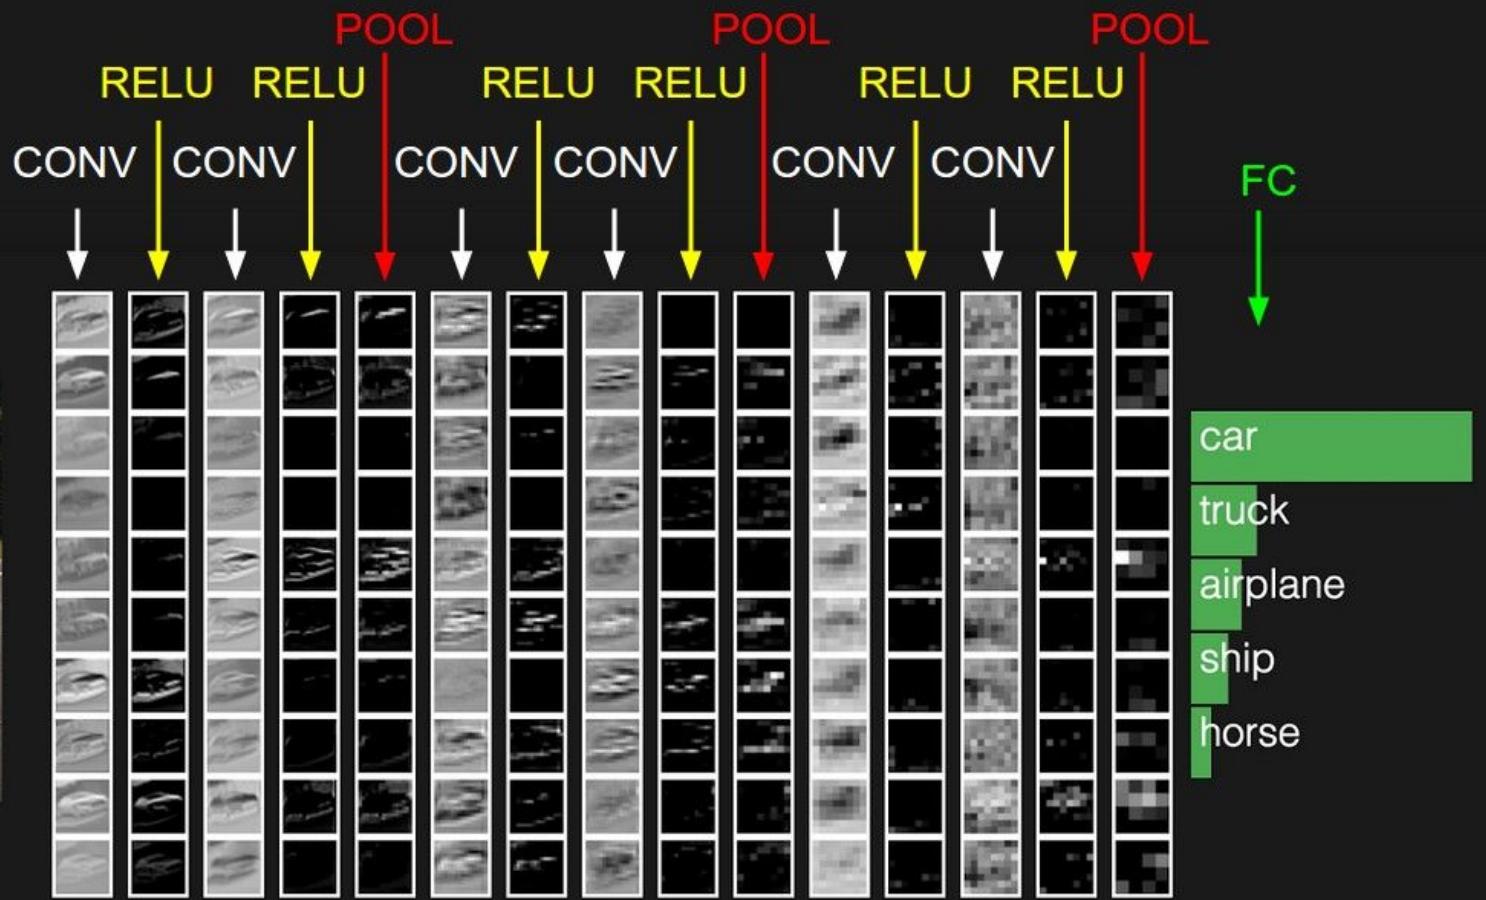
Previsão do tempo para Porto Alegre

Applying Convolutional Neural Networks Concepts To Hybrid NN Model For Speech Recognition.
Abdel-Hamid et al (2012)

The Convolutional Neural Network



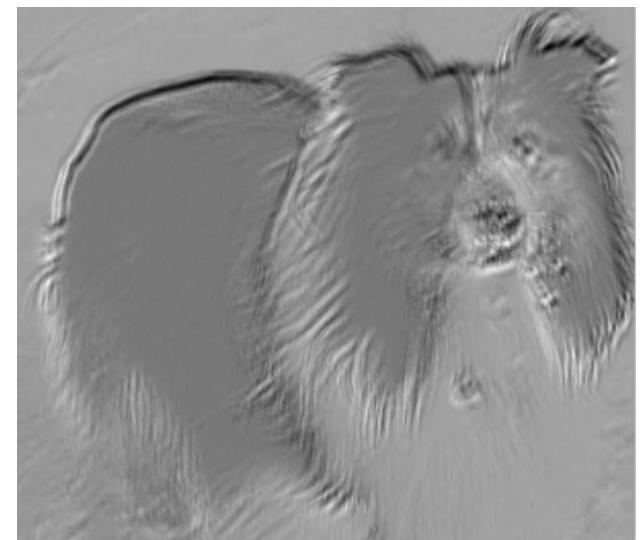
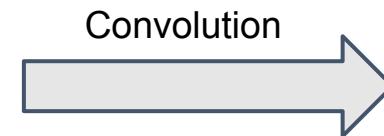
The Convolutional Neural Network



RELU: Rectified Linear Units

Fonte: <http://cs231n.stanford.edu/>

The Convolutional Neural Network



The Convolutional Neural Network

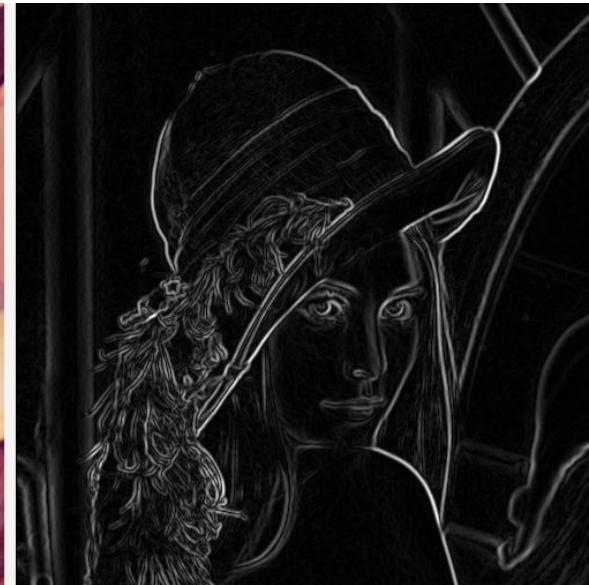
→ What is a convolution (*)?

$$G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * A \quad \text{e} \quad G_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * A$$

Image A



Resulting image: $G_x + G_y$



The Convolutional Neural Network

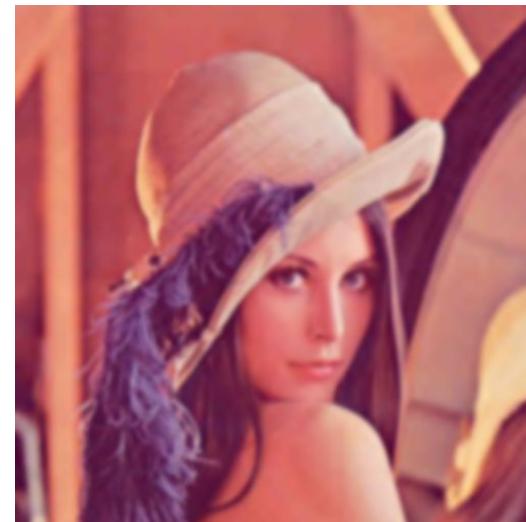
→ What is a convolution (*)?

$$G = \frac{1}{273} \cdot \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{bmatrix} * A$$

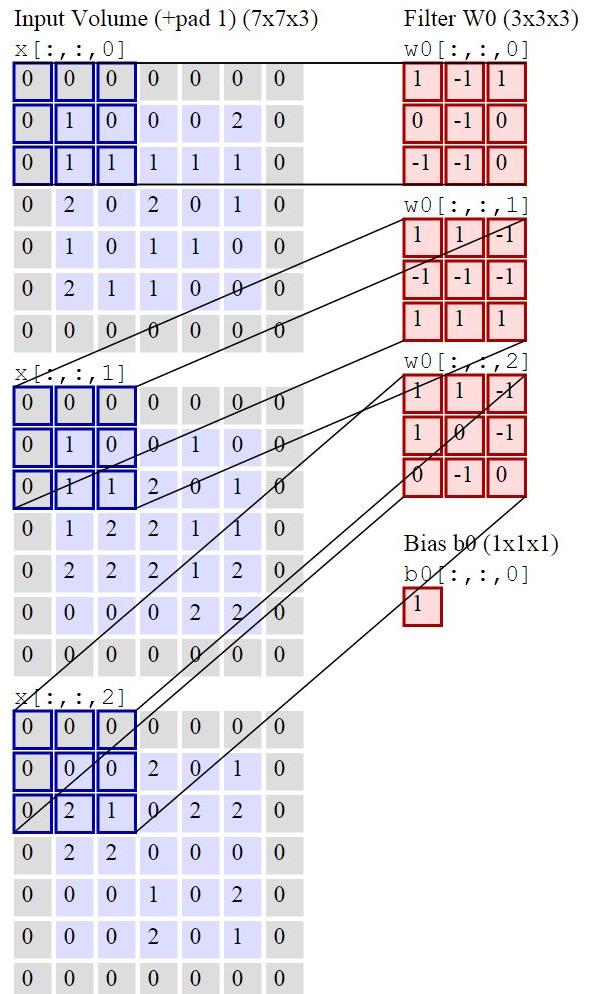
Image A



Resulting image G



The Convolutional Neural Network



Filter W0 (3x3x3)

$w0[:, :, 0]$

$w0[:, :, 1]$

$w0[:, :, 2]$

Bias $b1$ (1x1x1)

$b1[:, :, 0]$

0

Filter W1 (3x3x3)

$w1[:, :, 0]$

$w1[:, :, 1]$

$w1[:, :, 2]$

Bias $b1$ (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$\circ[:, :, 0]$

$\circ[:, :, 1]$

$\circ[:, :, 2]$

Input Image: 5x5x3

Padding: 1

Input Image (+pad): 7x7x3

Filters: 2

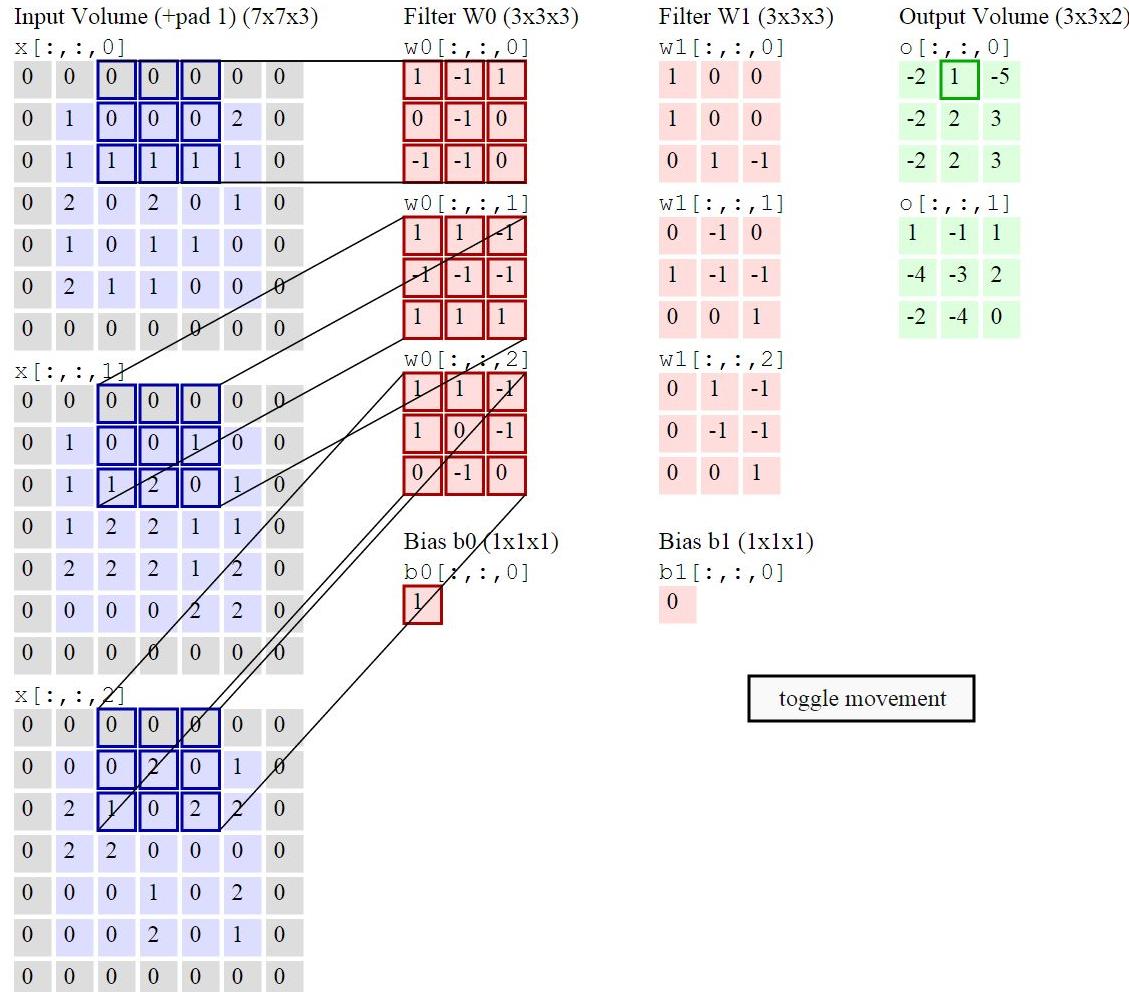
Filters Shape: 3x3

Stride: 2

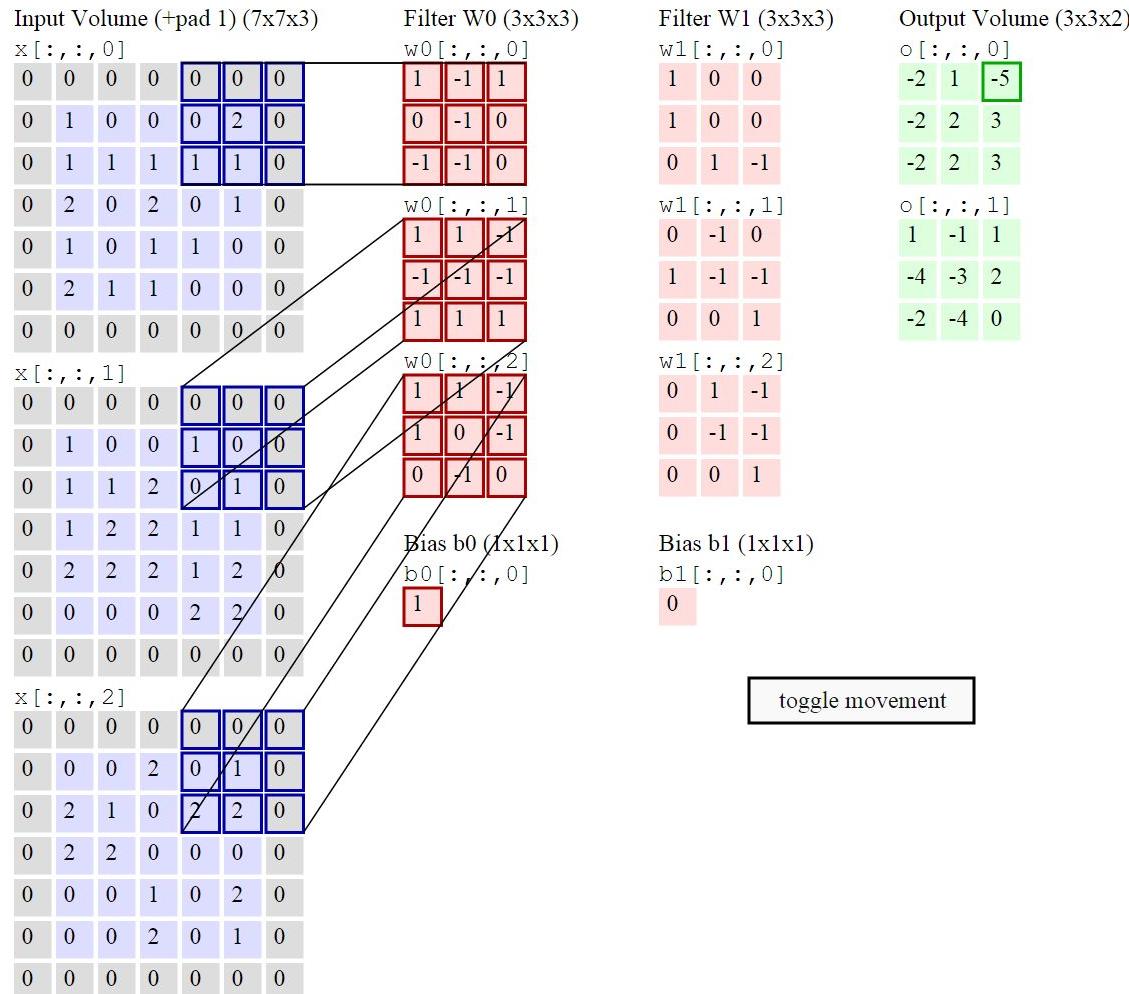
$$(5 + 1 * 2 - 3) / 2 + 1 = 3$$

toggle movement

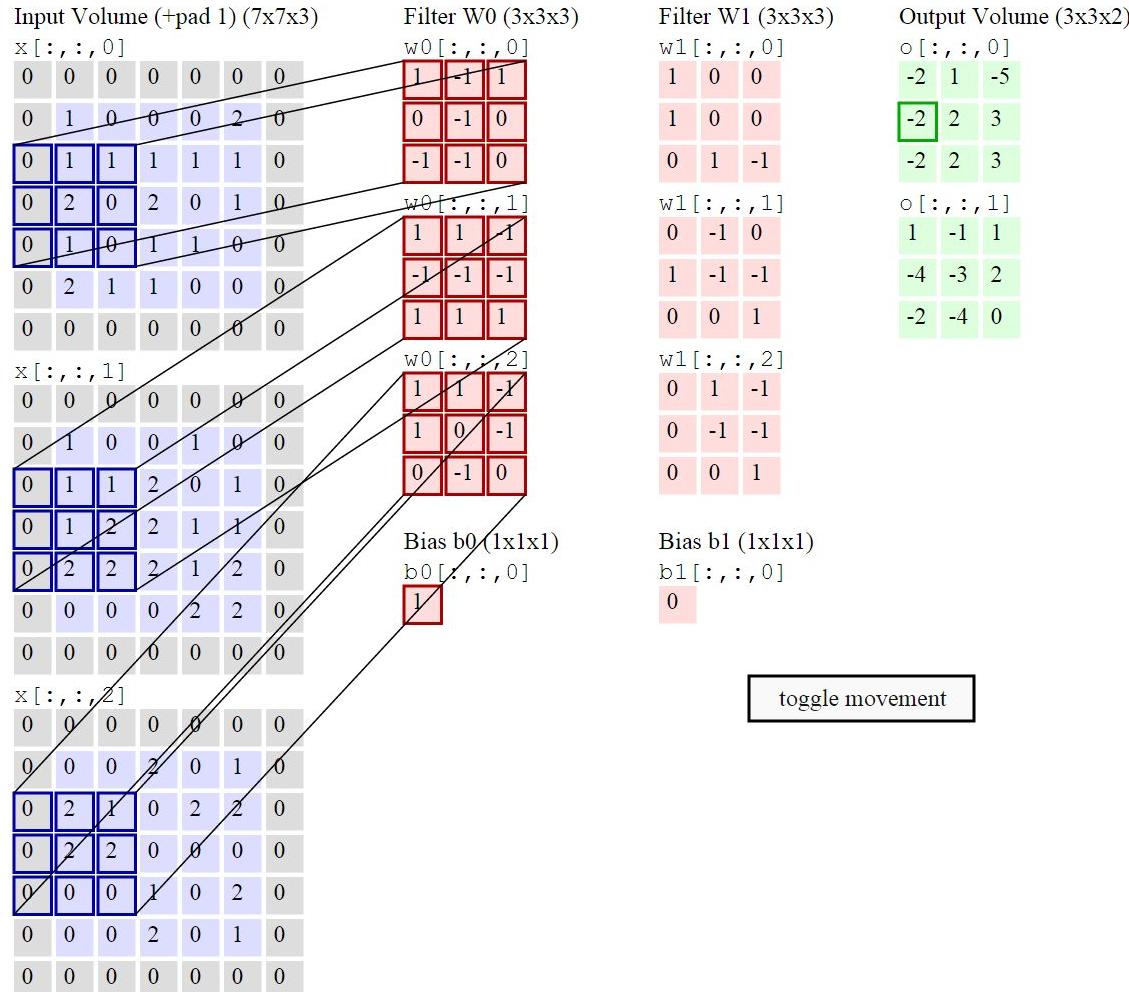
The Convolutional Neural Network



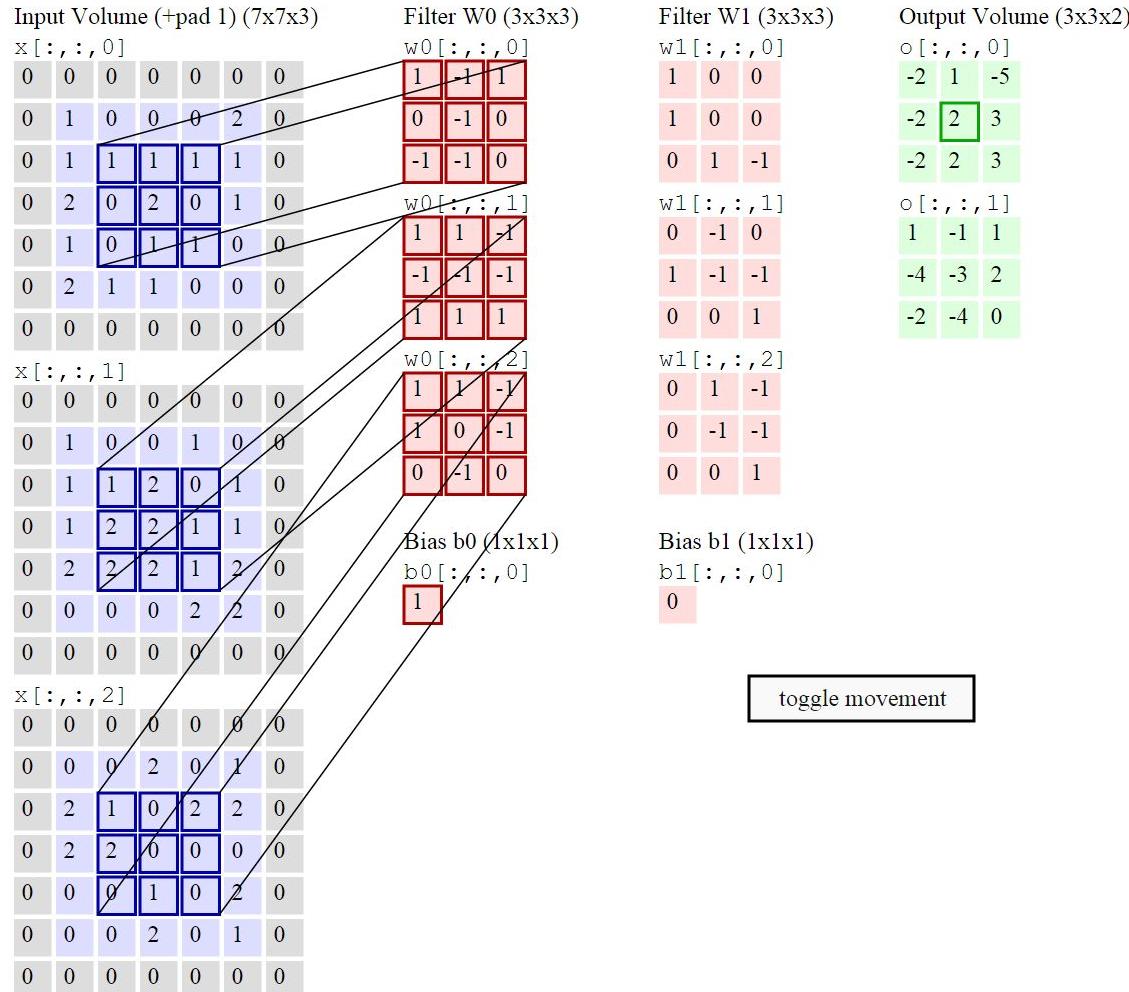
The Convolutional Neural Network



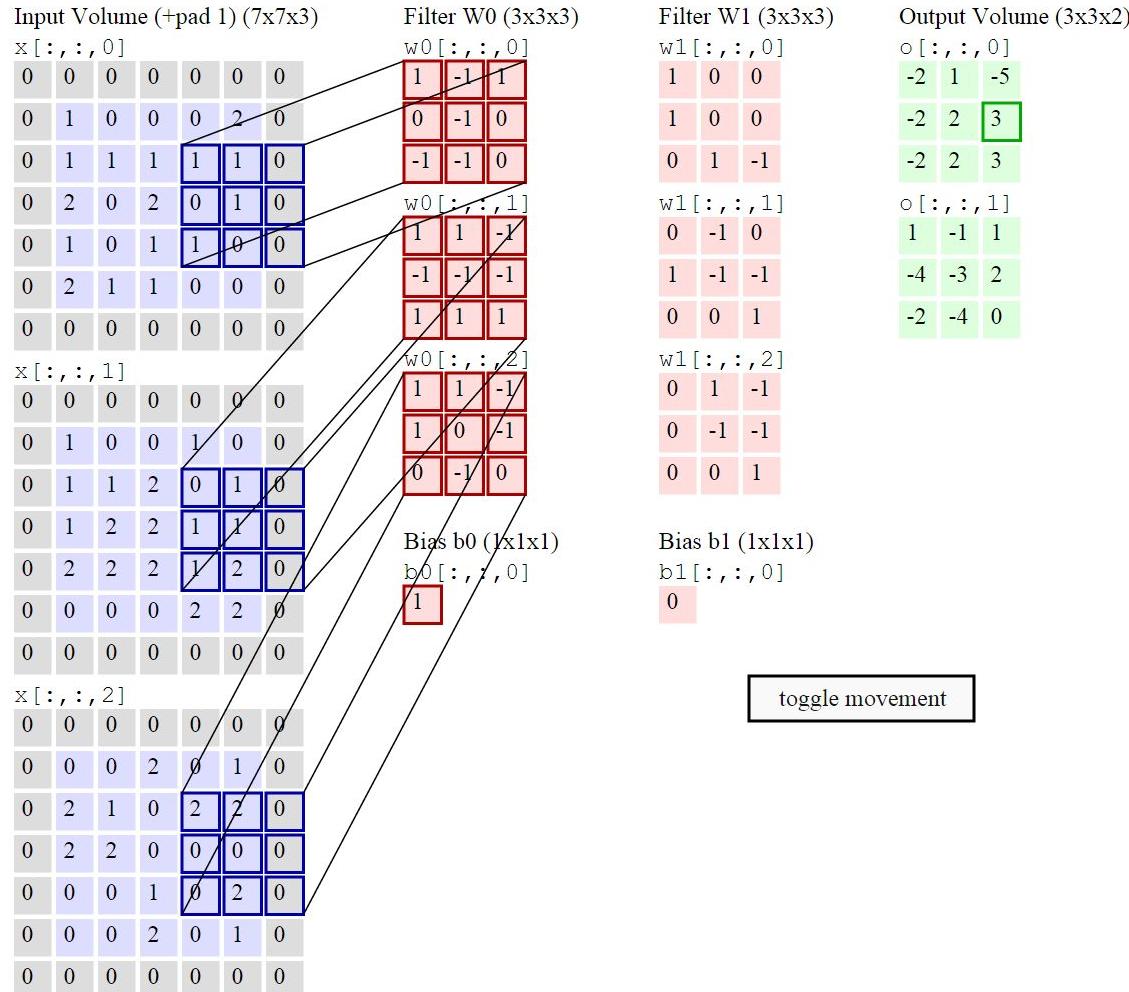
The Convolutional Neural Network



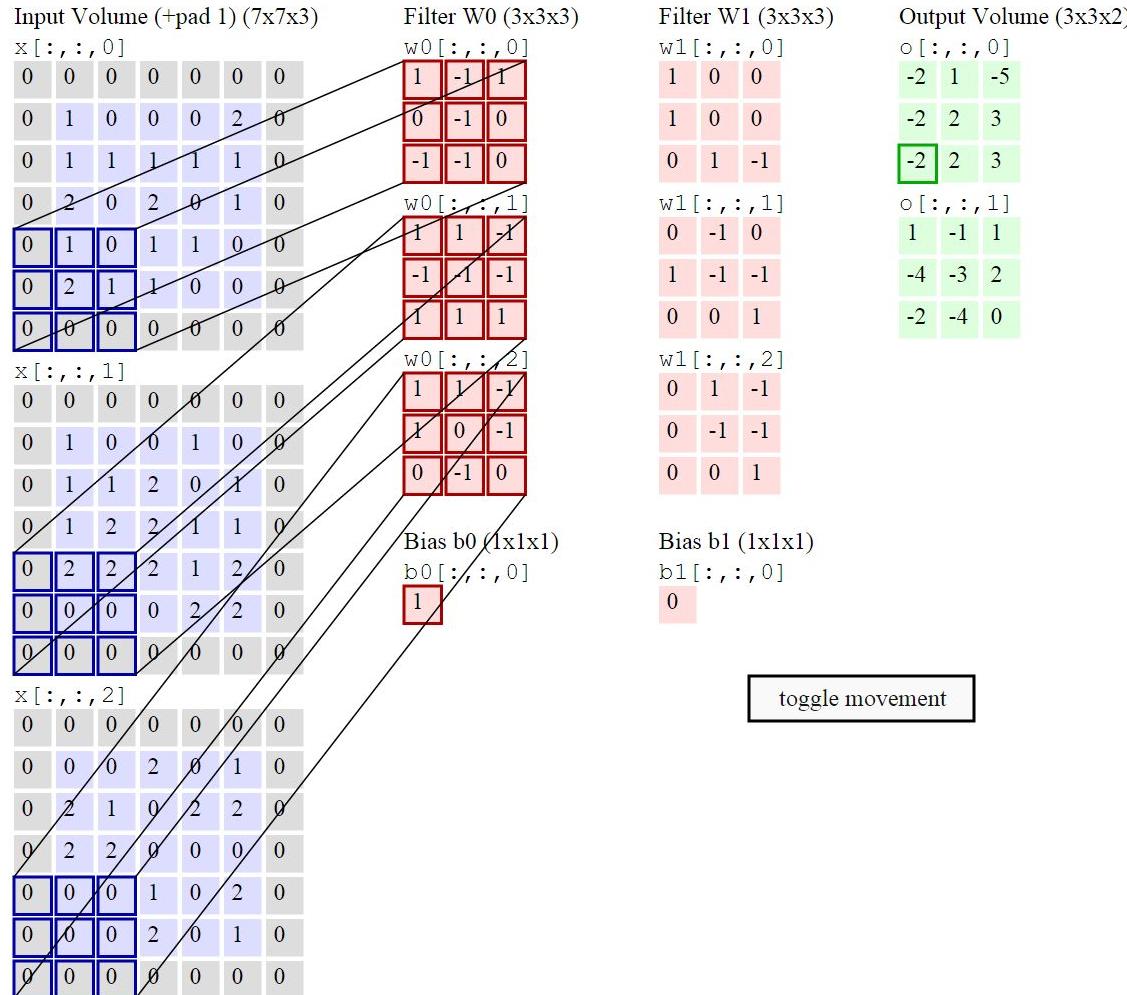
The Convolutional Neural Network



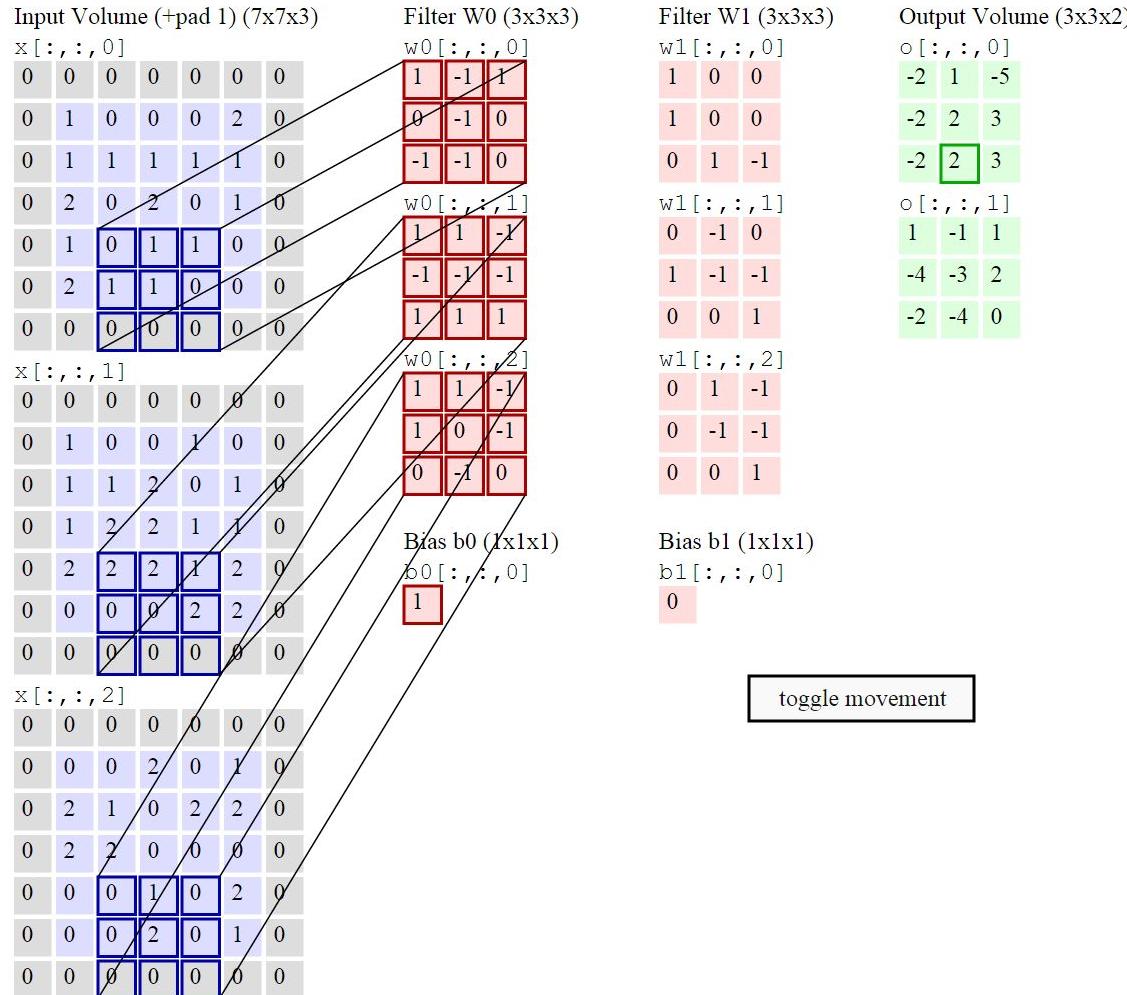
The Convolutional Neural Network



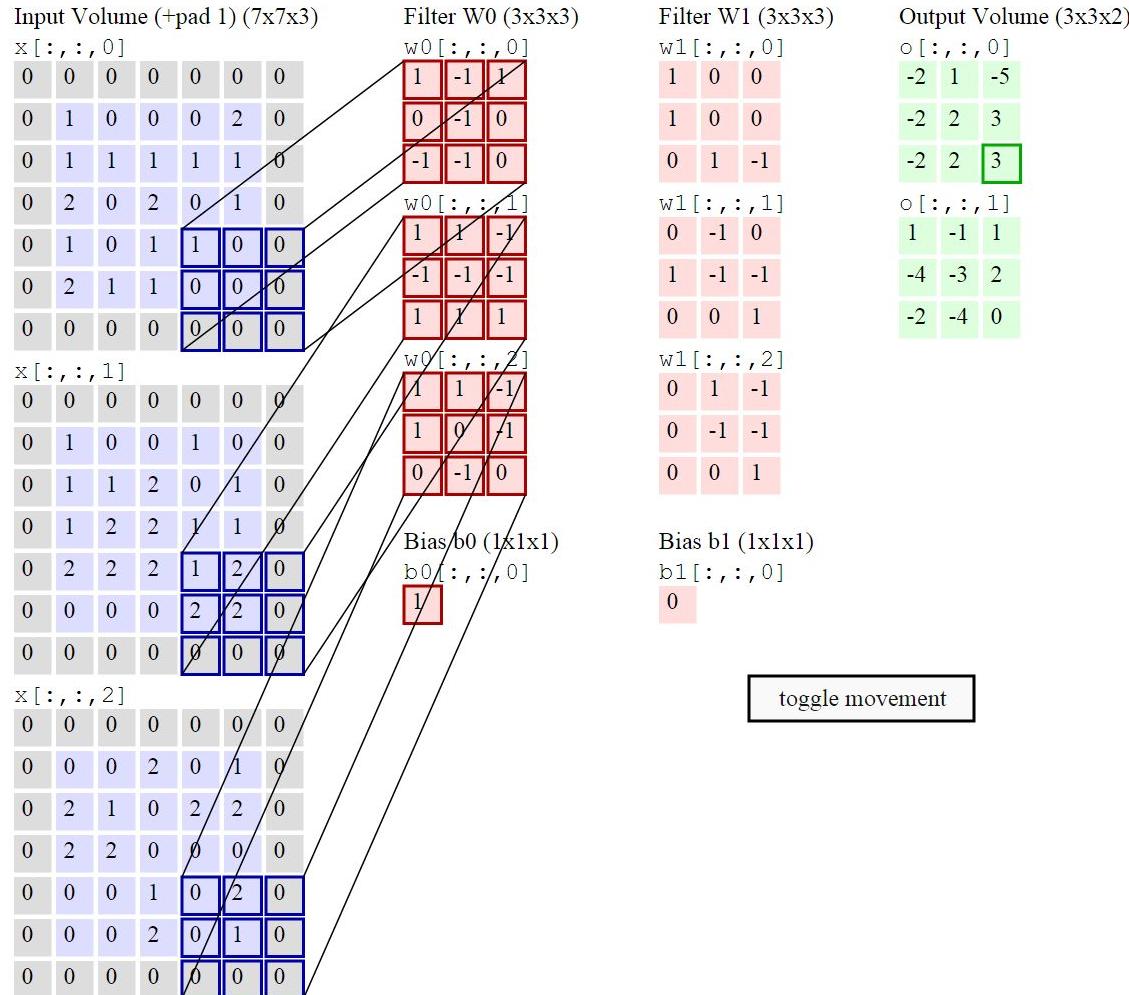
The Convolutional Neural Network



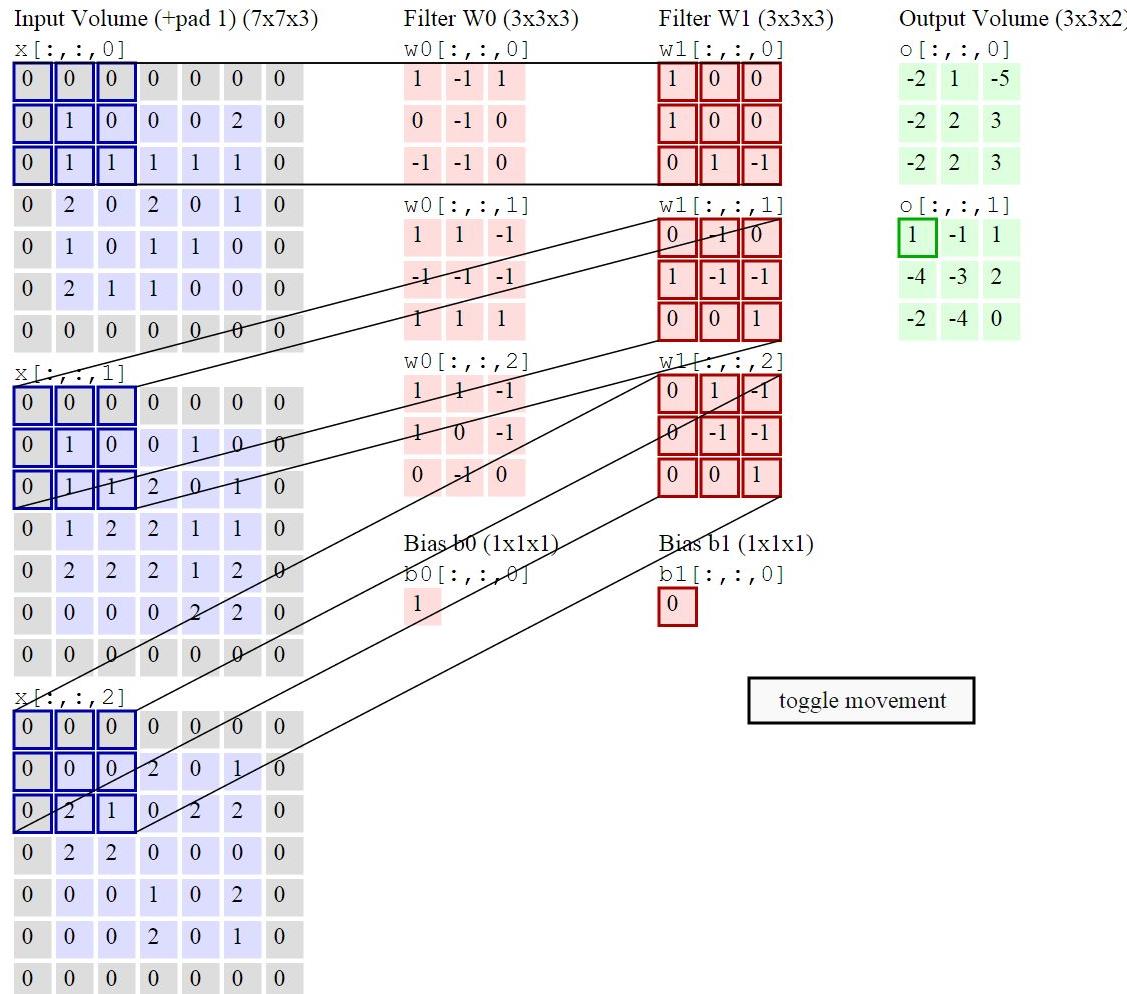
The Convolutional Neural Network



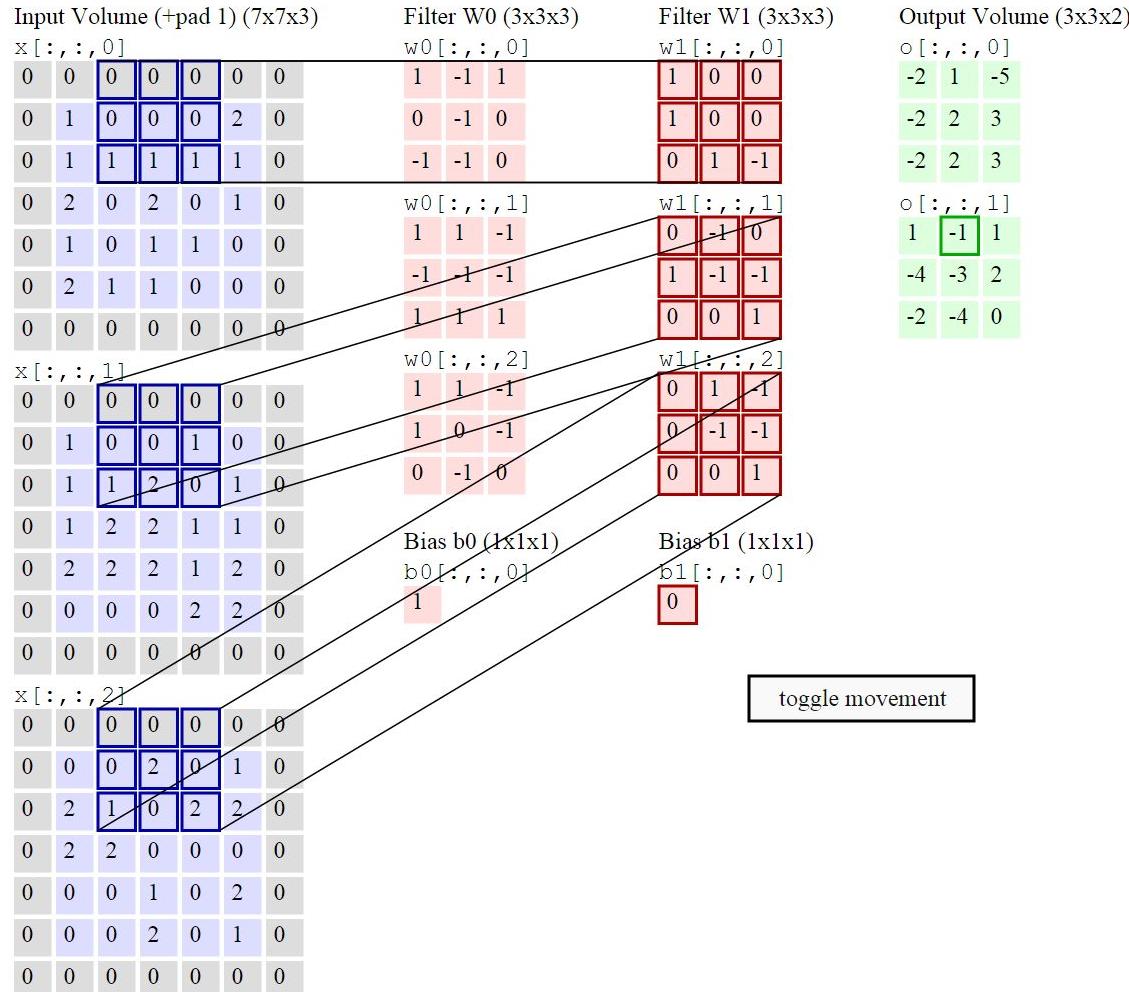
The Convolutional Neural Network



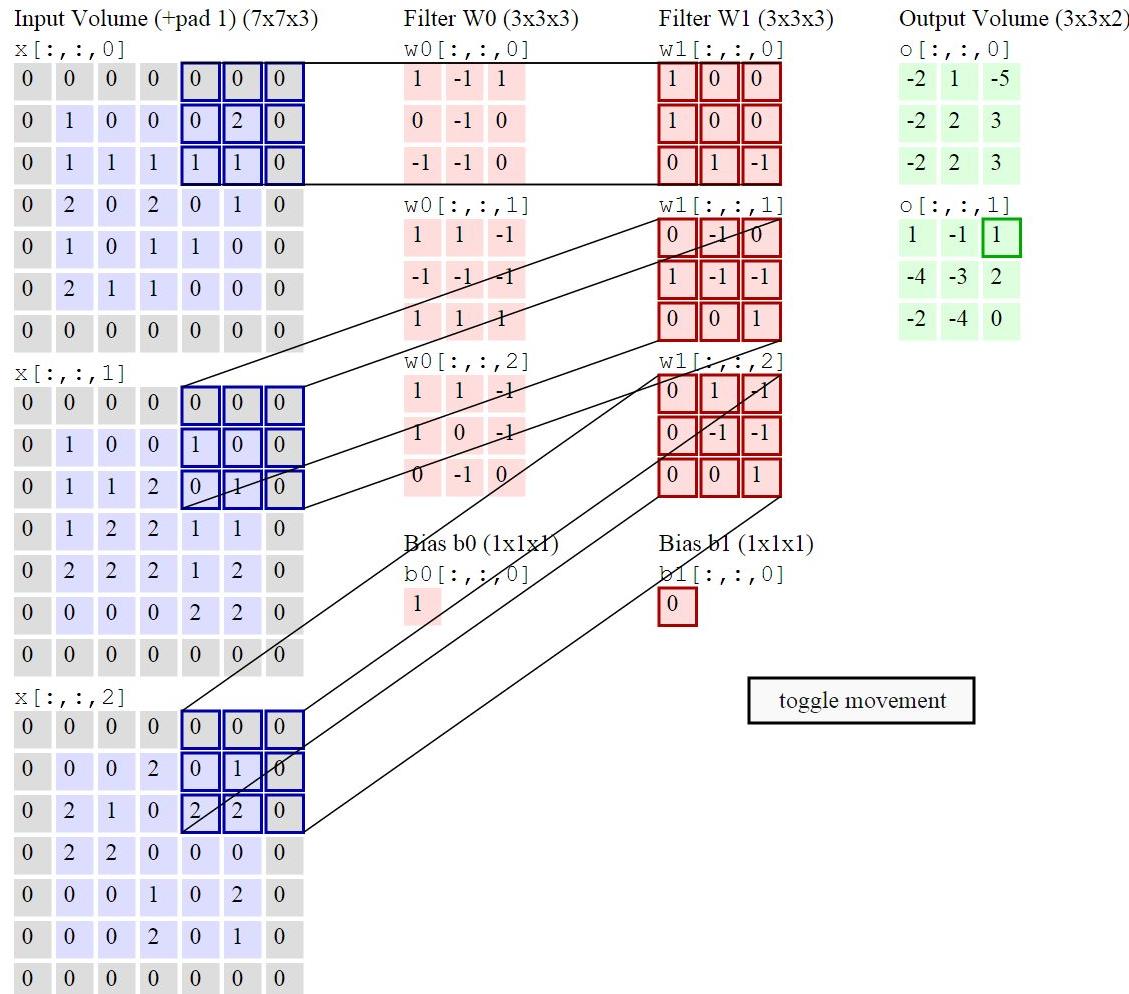
The Convolutional Neural Network



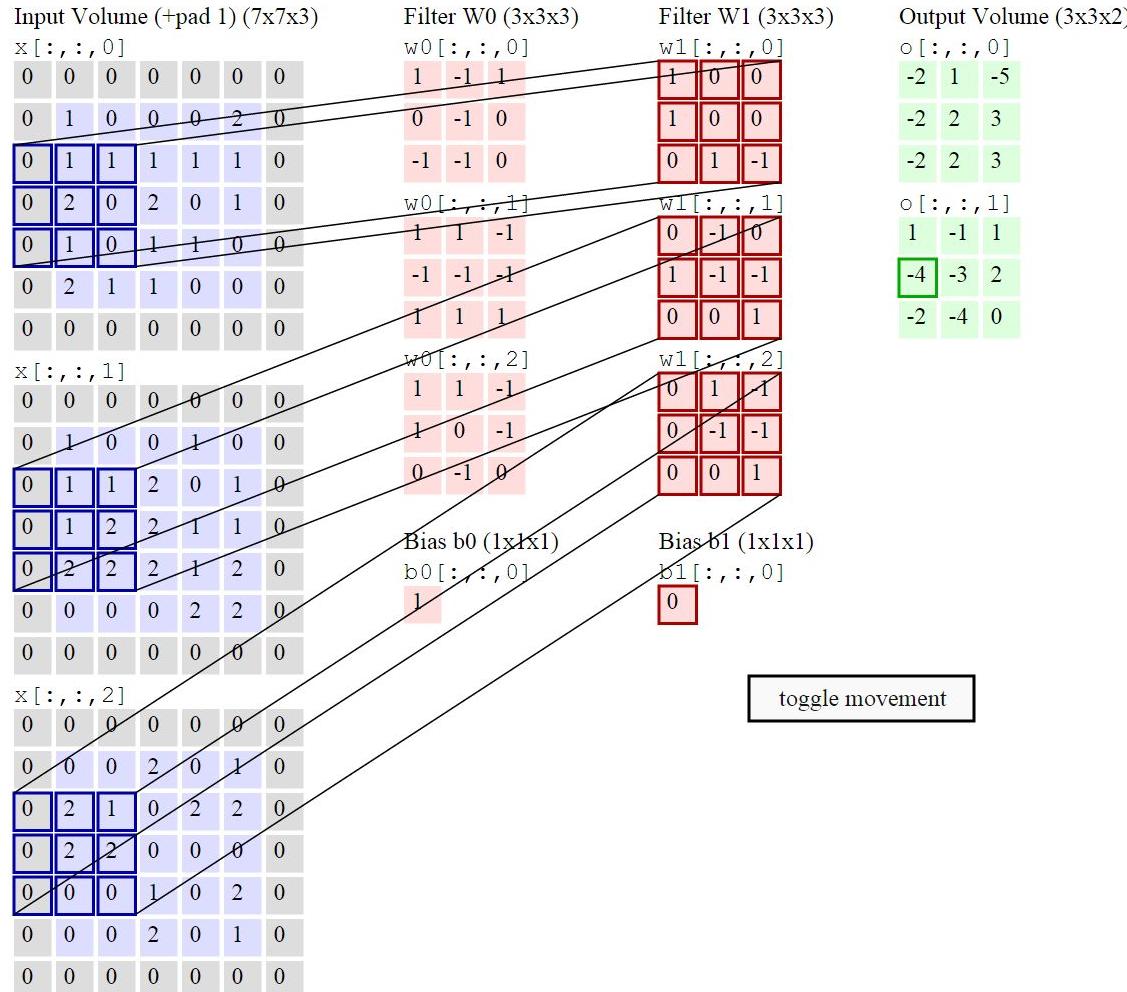
The Convolutional Neural Network



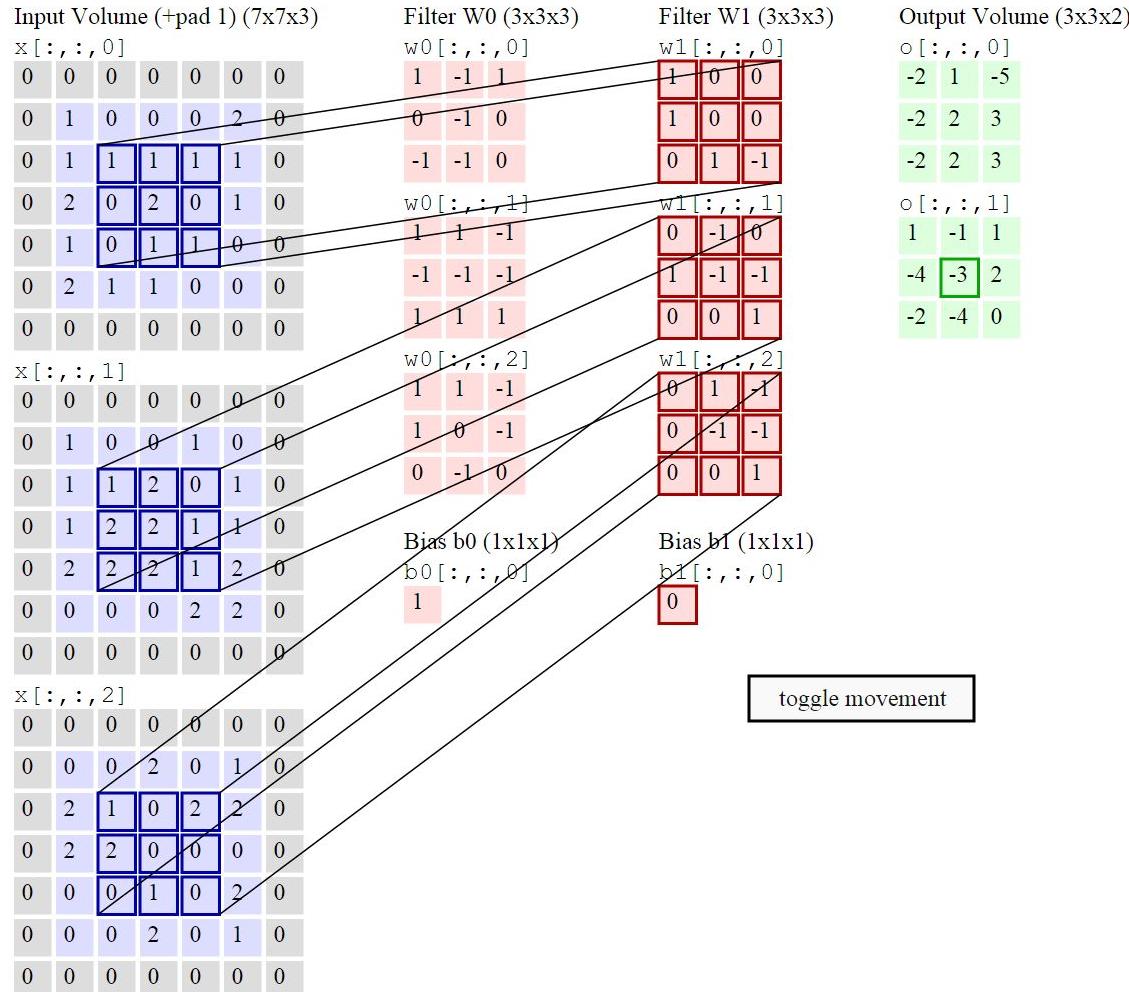
The Convolutional Neural Network



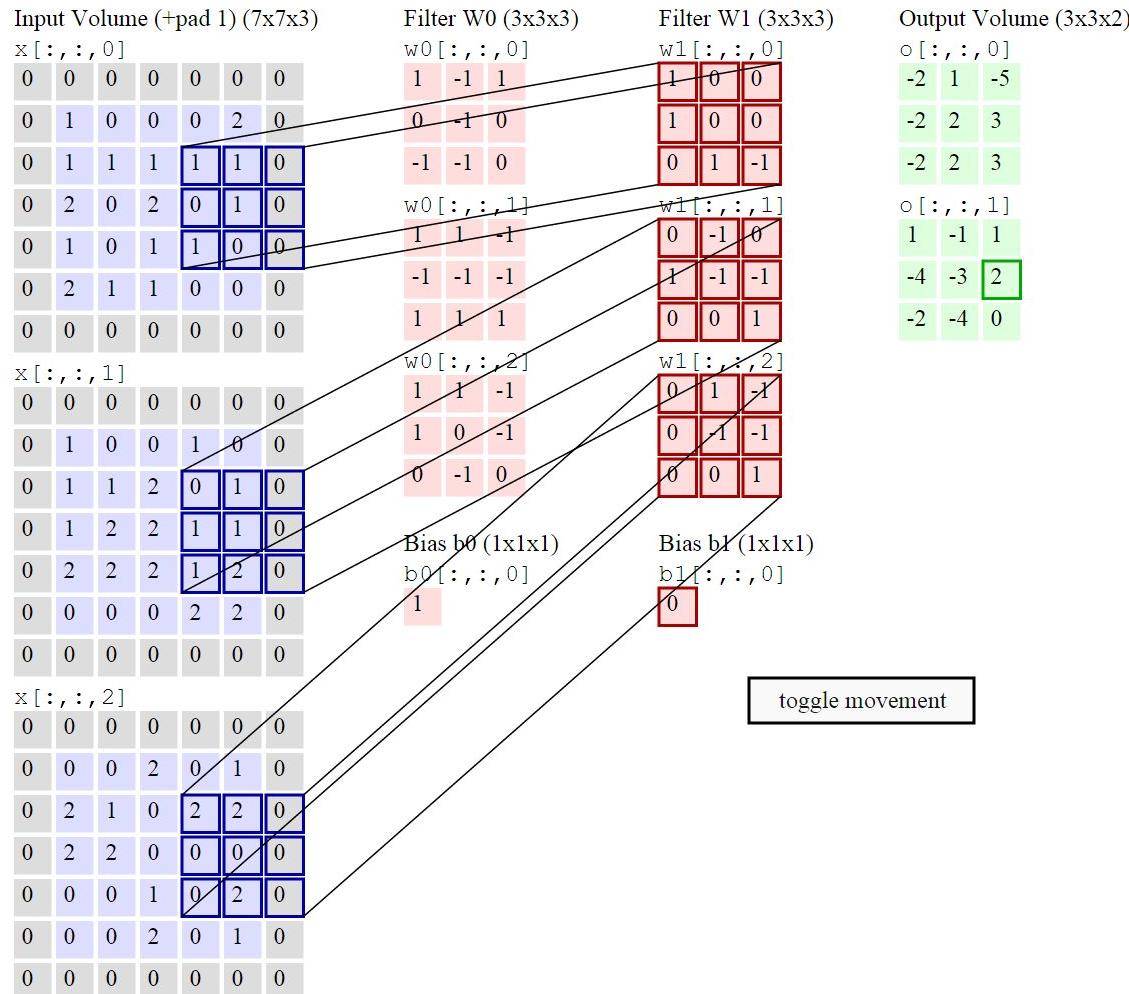
The Convolutional Neural Network



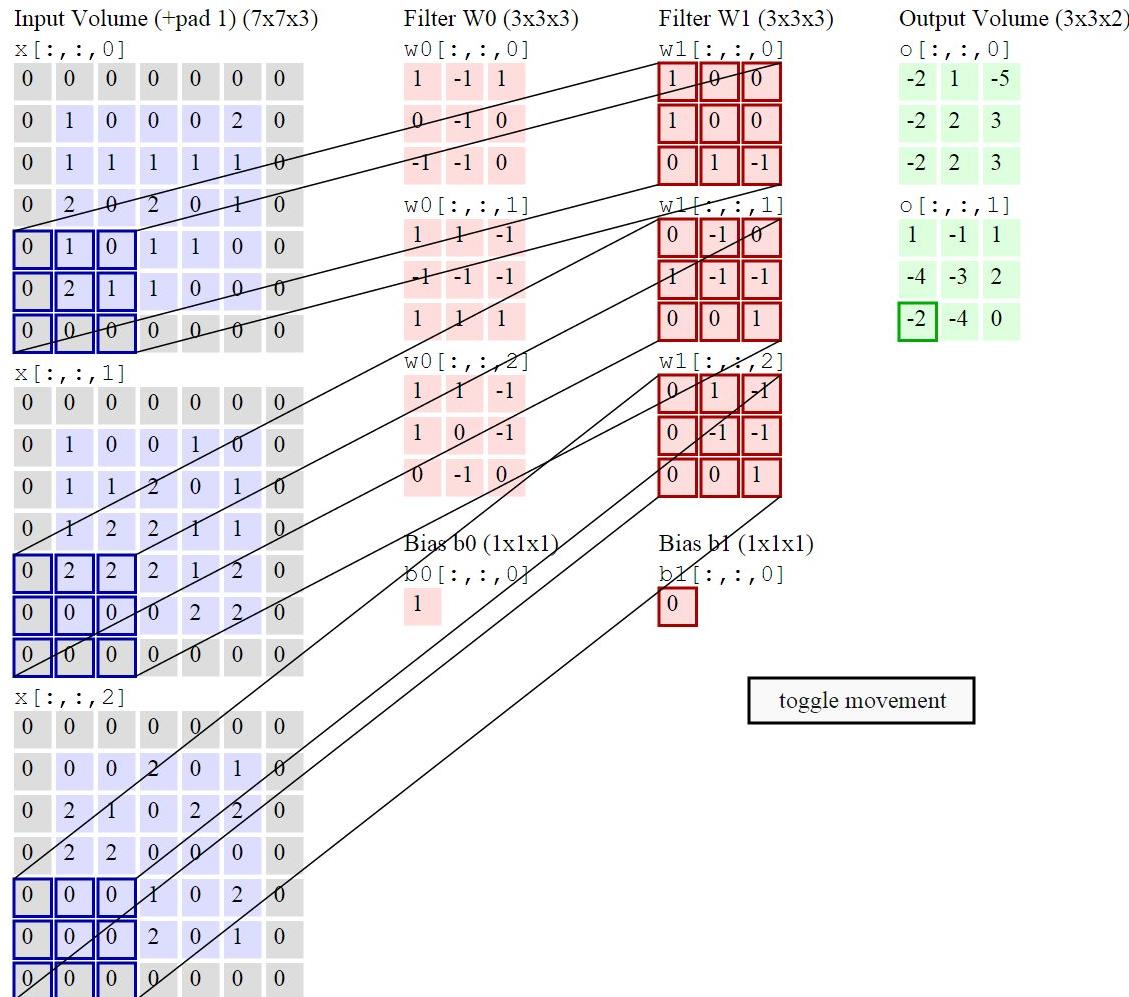
The Convolutional Neural Network



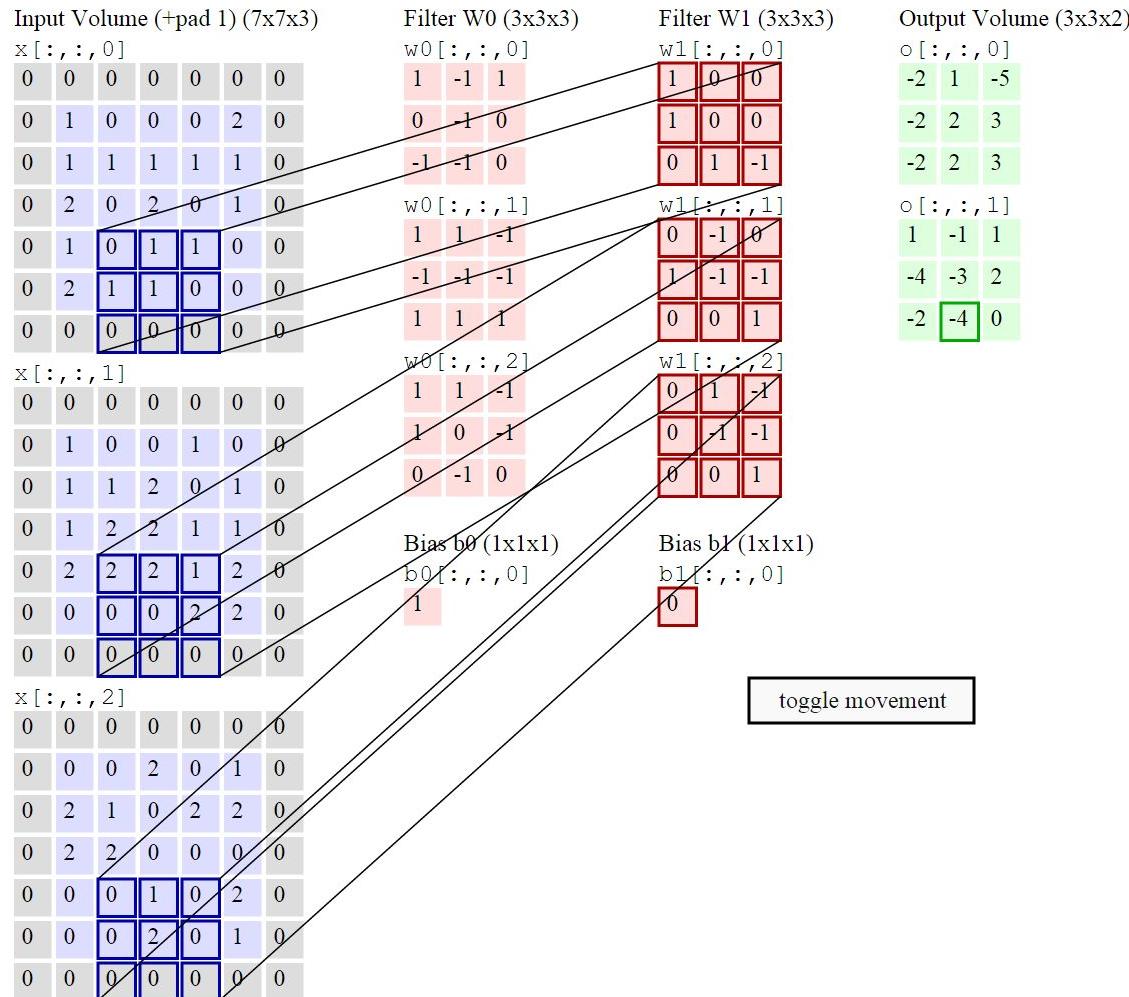
The Convolutional Neural Network



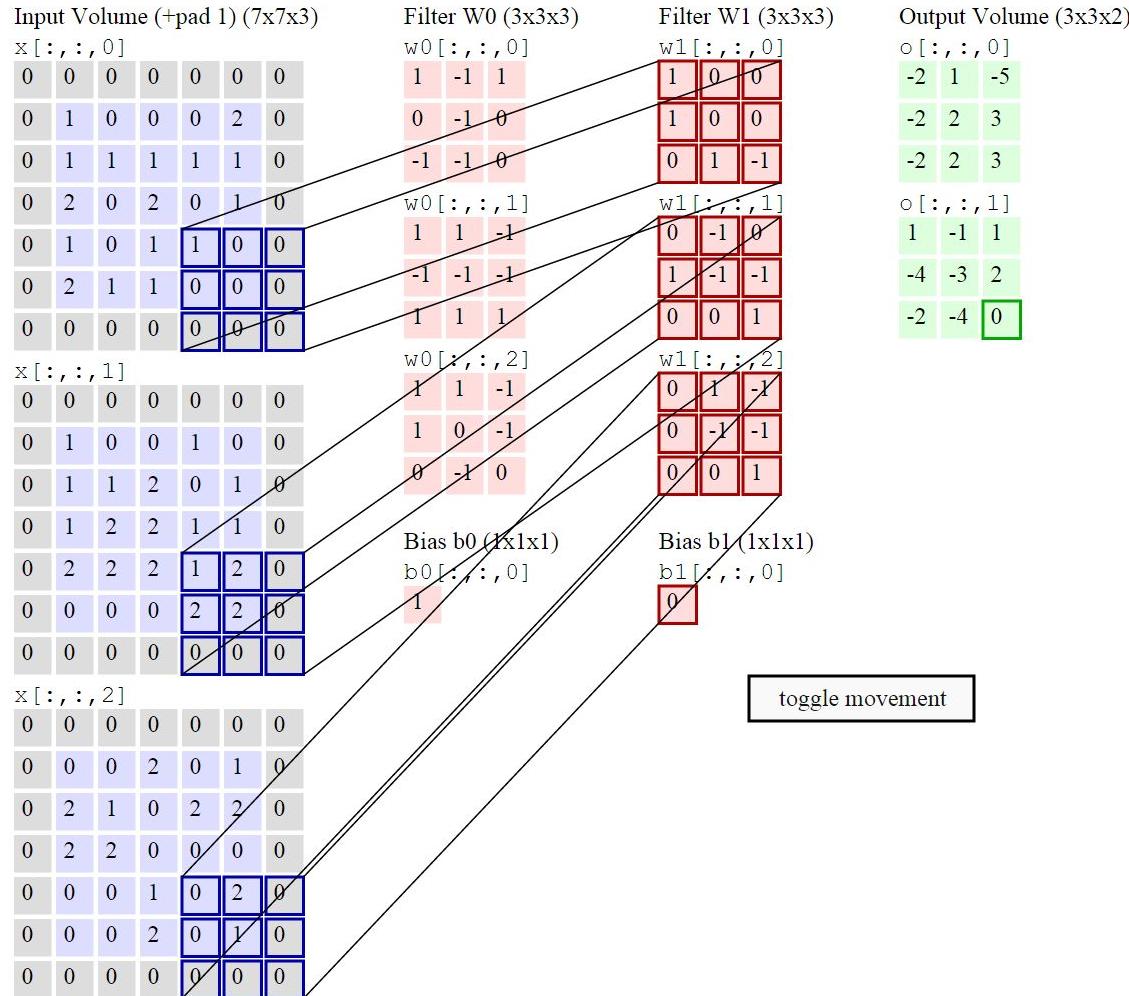
The Convolutional Neural Network



The Convolutional Neural Network



The Convolutional Neural Network



Convolving a Dog

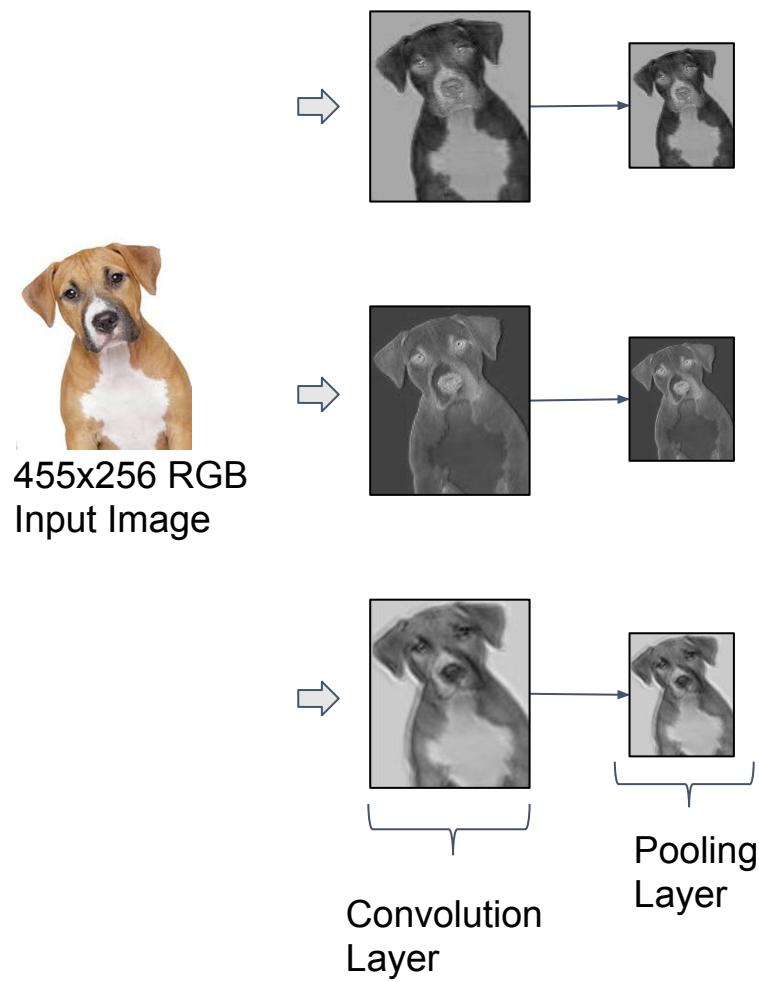


455x256 RGB
Input Image



Convolution
Layer

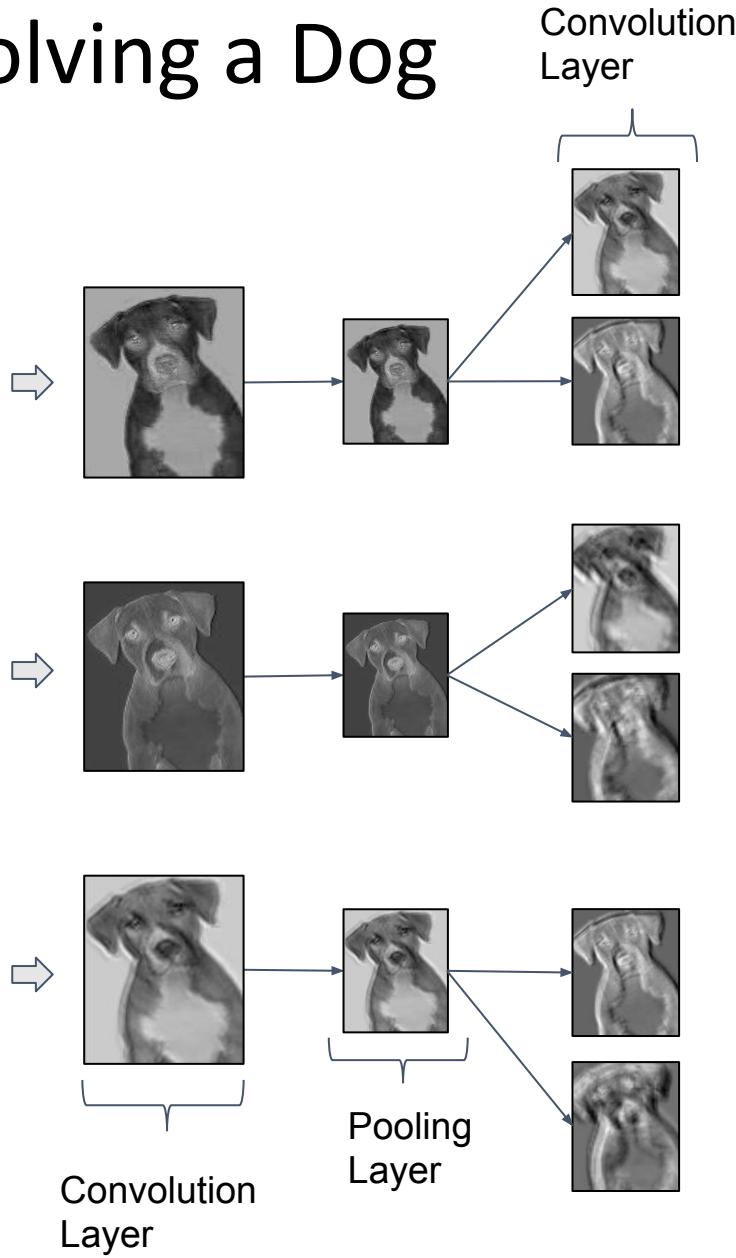
Convolving a Dog



Convolving a Dog

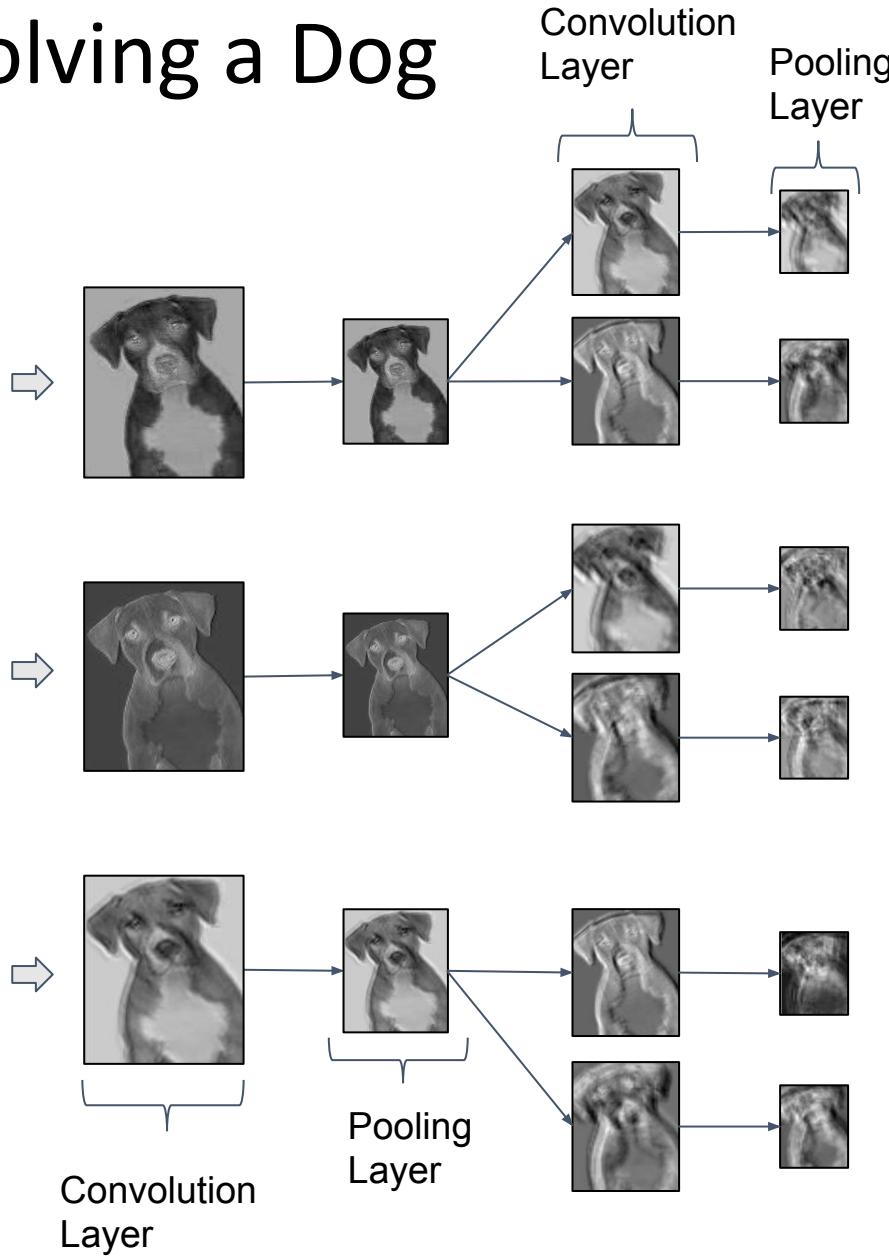


455x256 RGB
Input Image



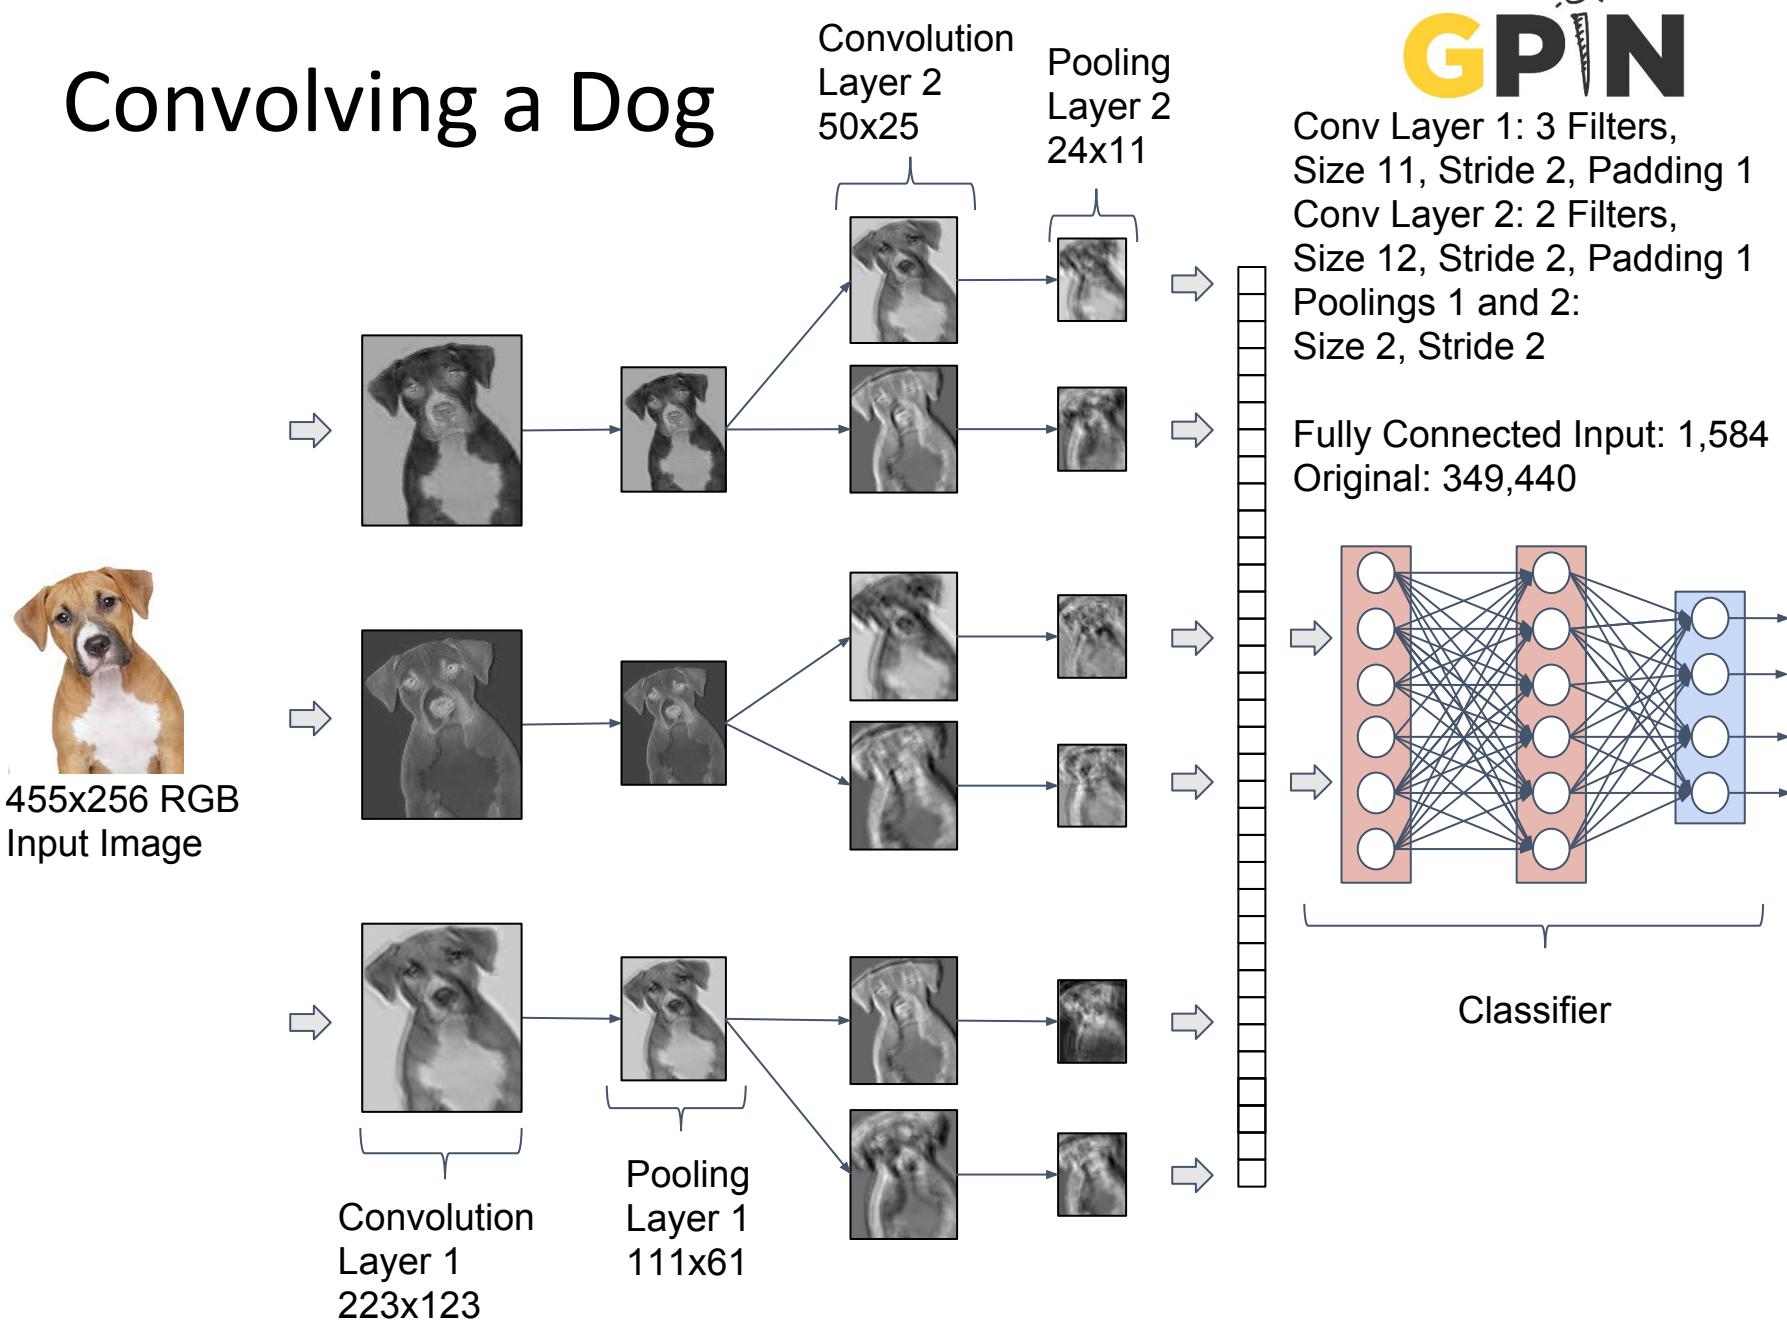
Convolving a Dog


455x256 RGB
Input Image



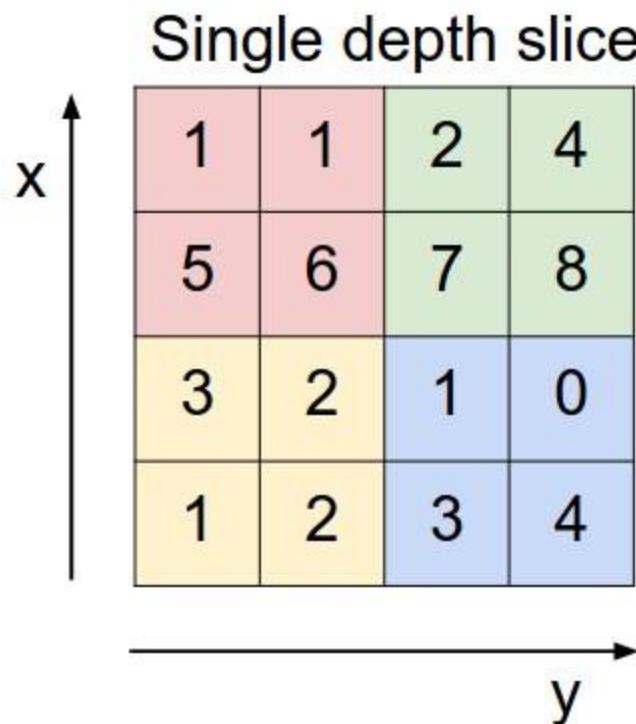


Convolving a Dog

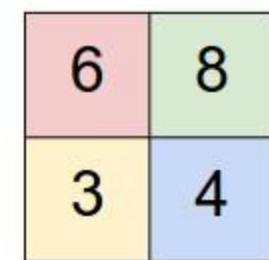


The Convolutional Neural Network

Pooling layer



max pool with 2x2 filters
and stride 2

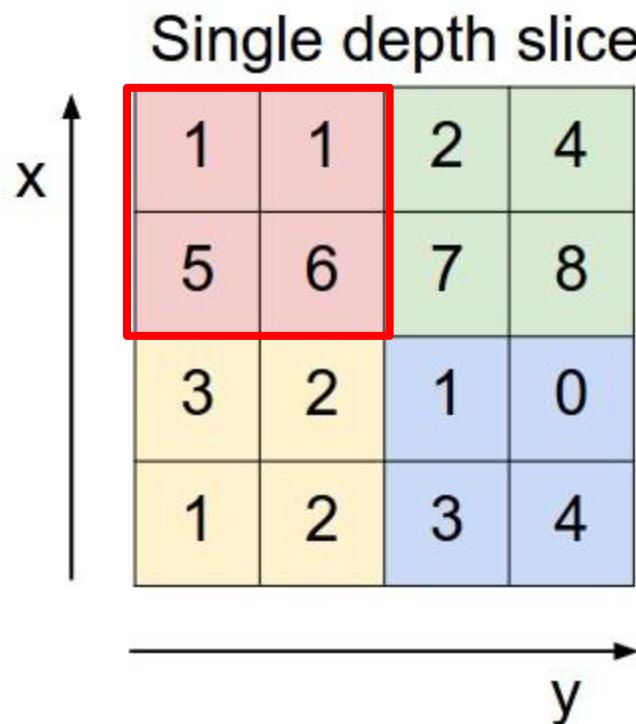


A 2x2 grid representing the output of the max pooling layer. It contains four values: 6 (top-left), 8 (top-right), 3 (bottom-left), and 4 (bottom-right). The cells are colored according to the input matrix: pink for the top-left cell, green for the top-right, yellow for the bottom-left, and blue for the bottom-right.

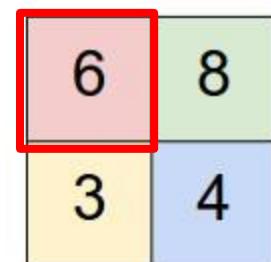
6	8
3	4

The Convolutional Neural Network

Pooling layer

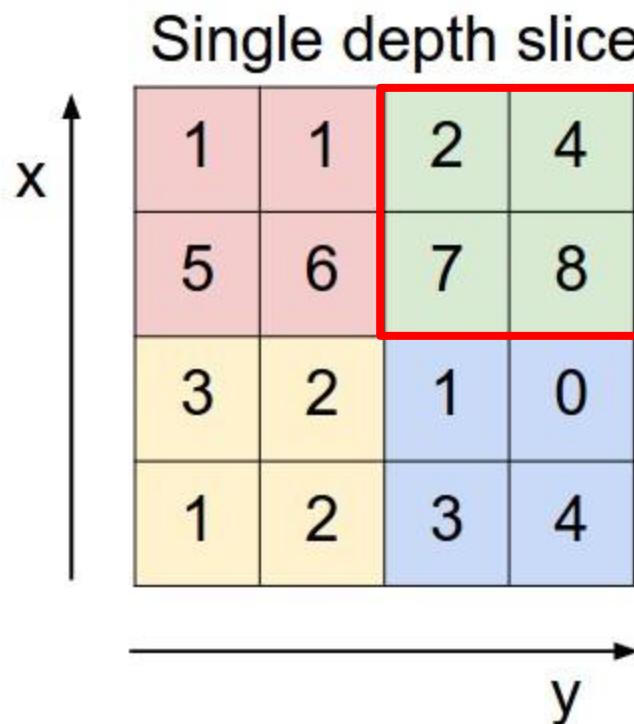


max pool with 2x2 filters
and stride 2

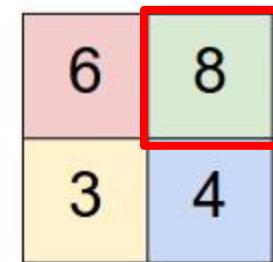


The Convolutional Neural Network

Pooling layer

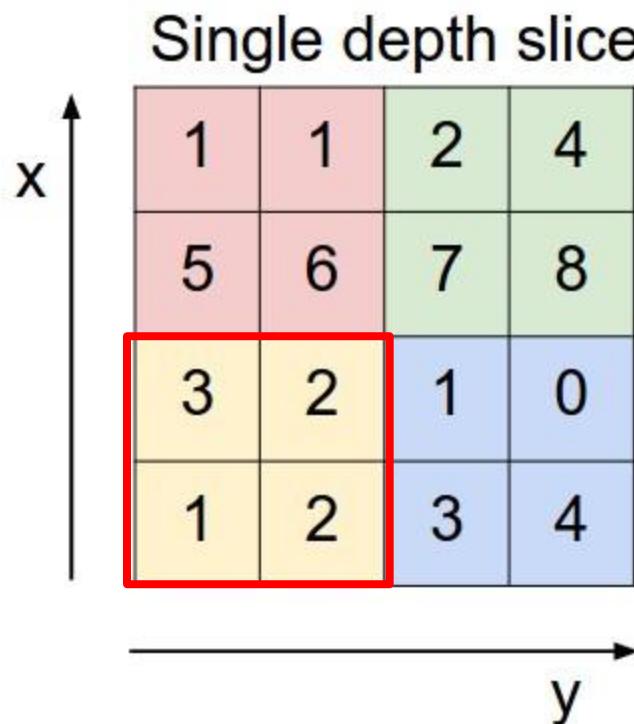


max pool with 2x2 filters
and stride 2

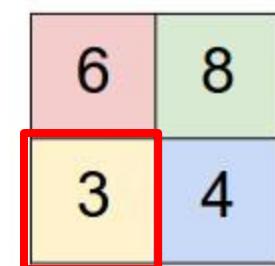


The Convolutional Neural Network

Pooling layer

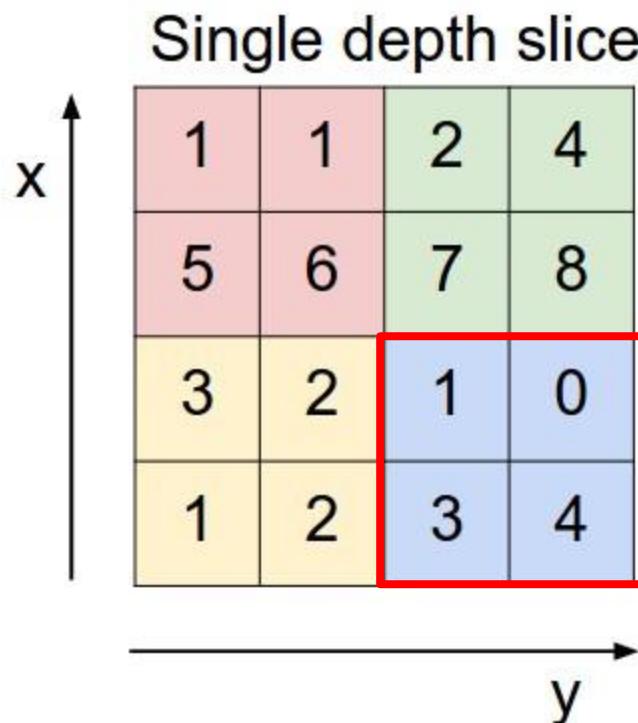


max pool with 2x2 filters
and stride 2

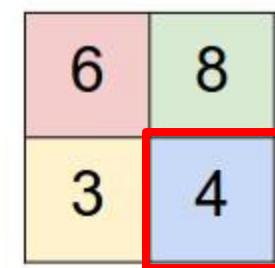


The Convolutional Neural Network

Pooling layer



max pool with 2x2 filters
and stride 2



Input: 4x4x1
Pooling Size: 2
Stride: 2

$$(4 - 2) / 2 + 1 = 2$$

Deep Learning Frameworks

Where the magic turns to be!

What is a deep learning framework?

It is a simple toolbox that implements neural networks concepts in a loose coupled approach. Most of this concepts is presented as building blocks, commonly known as layers.

A deep neural network is a set of different layers.

What is a deep learning framework?

Deep learning frameworks also implement algorithms for training neural networks based on backpropagation, including optimizers, parameter initializers, and loss functions.

Most popular frameworks

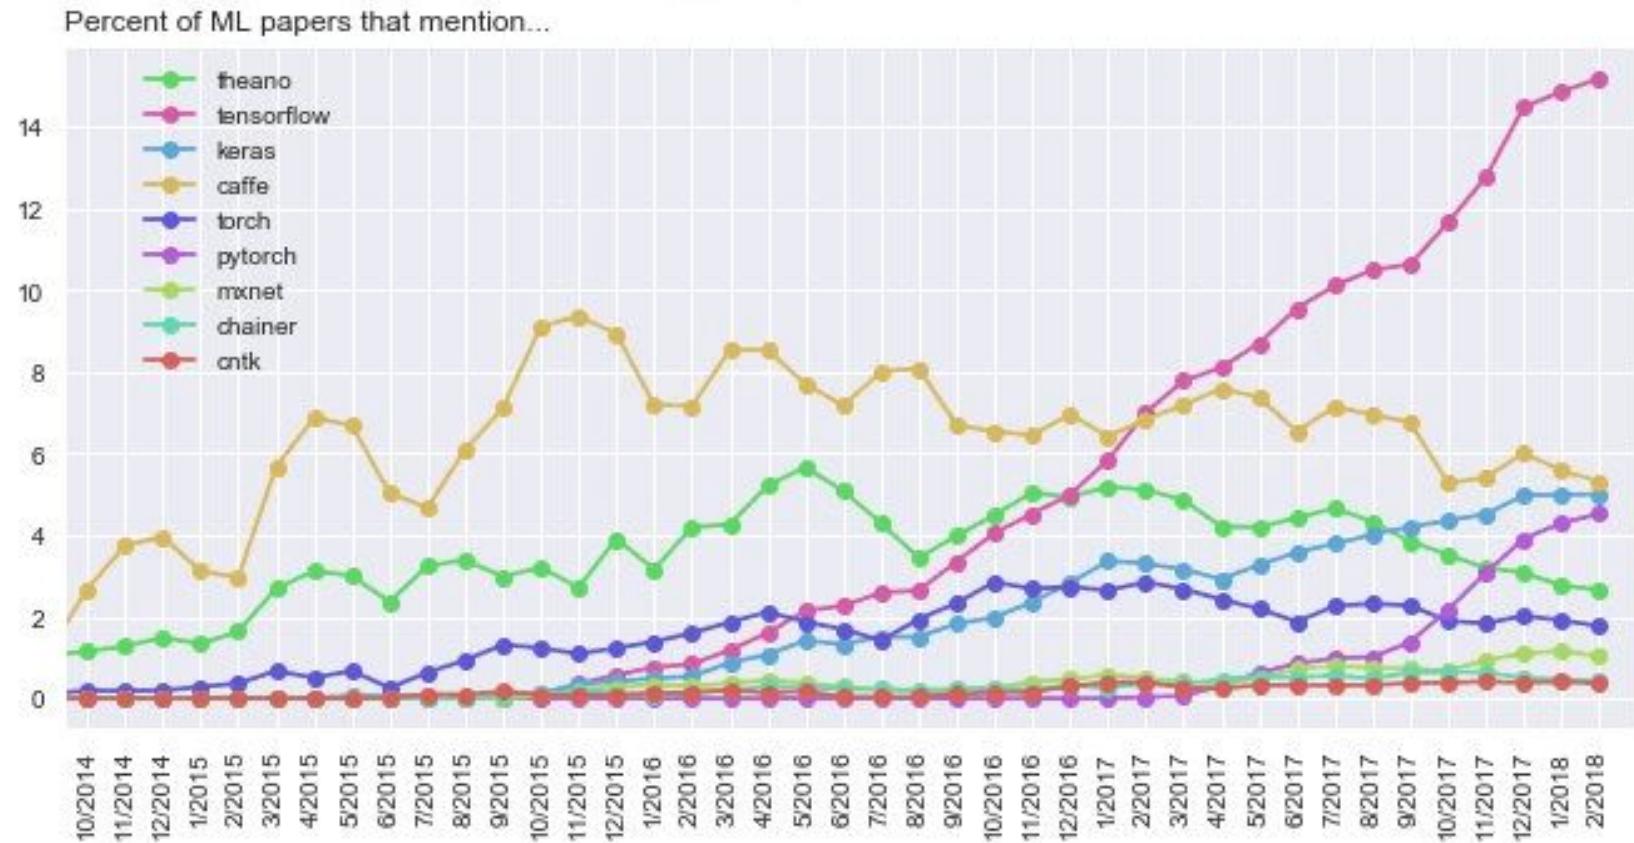
- PyTorch
- Tensorflow
- Keras
- Caffe

Most popular frameworks

So, which framework is the best choice?

There is no answer! But...

Most popular frameworks



Some practice...

To do some practice in our discipline, we will use a virtual machine with some pre-installed frameworks.

This approach is not ideal. We will use this virtual machine due to the lack of specific resources for deep learning, such as GPUs and high performance machines.

Some practice...

Let's take a look at what we have in
our sandbox.

Assignment 1 - Car Classifier

Our practice in the sandbox used several pre-trained models to classify contents. One of these models classifies American cars, however, probably this same model will have a fragile performance classifying Brazilian cars.

Assignment 1 - Car Classifier

Our practice in the sandbox used several pre-trained models to classify contents. One of these models classifies American cars, however, probably this same model will have a fragile performance classifying Brazilian cars.

Your challenge is to build a classifier capable of automatically identifying 5 models of Brazilian cars. To do this, use the parking lot of PUCRS and your mobile camera to take photos of different cars belonging to these 5 models. Construct your classifier using neural network features, but remember, you can use any machine learning concepts already seen in the course. Keep in mind that training a new deep model in our sandbox will not be a good idea.

Assignment 1 - Car Classifier

Important dates:

01/12: follow-up

08/12: follow-up

15/12: delivery date

You should deliver your source code, your dataset, and the best model you've gotten in your experiments.

Have Fun!